

Where We Are

Machine Learning Systems

2012 - Now

Big Data

2010 - Now

Cloud

2000 - 2016

Foundations of Data Systems

1980 - 2000



Logistics

- Exam date:
 - Final Exam date (tentative): **Friday, March 22, 8 - 11 am, PT**
 - **Decision: In-person Exam**
- Next week:
 - TA will hold multiple hours of Exam review
 - Pay attention to Piazza announcement about scheduling
 - Make sure you attend and get important secret sauces 😊
 - TAs and I will all be available for OH by appointment to help you on exam and wrapping the course!

ML System history

- ML Systems evolve as more and more ML components (models/optimization algorithms) are unified

Ad-hoc: diverse model family,
optimization algos, and data

Opt algo: iterative-convergent

Model family: neural nets

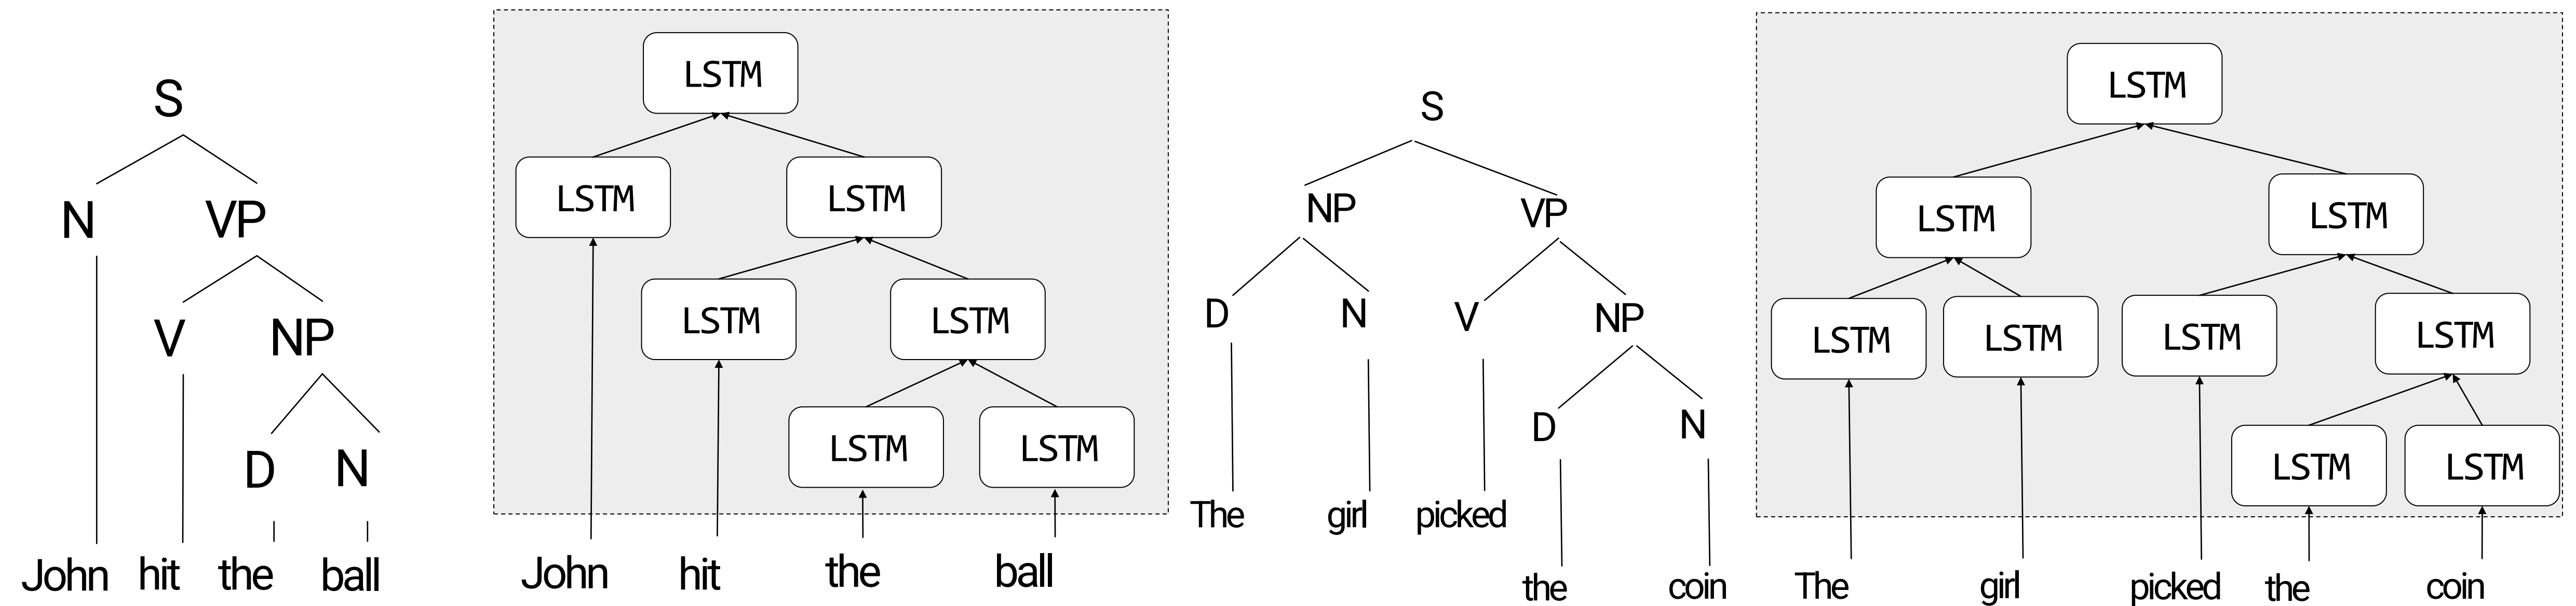
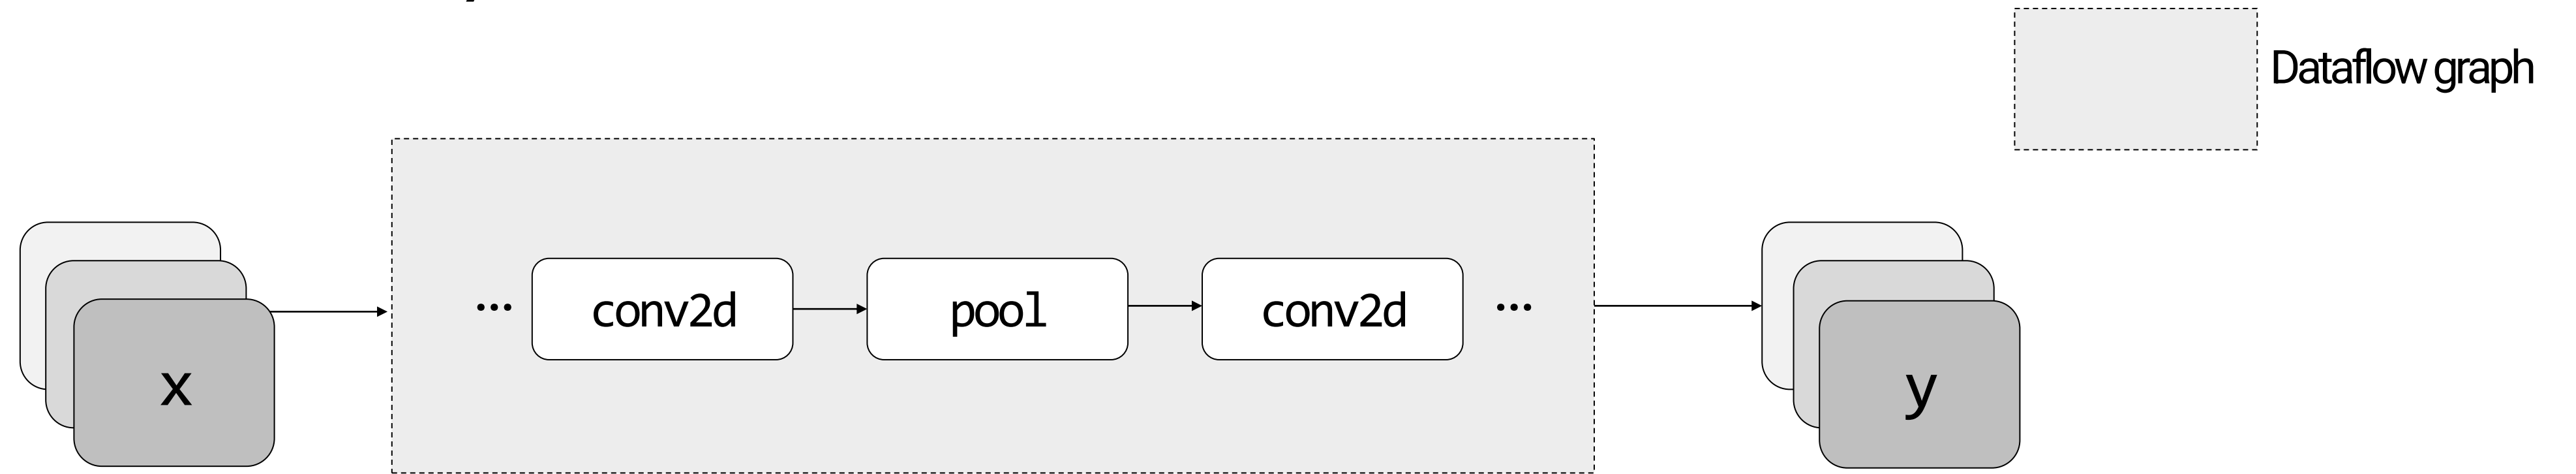
Model:
CNNs/transformers/GNNs

LLMs: transformer
decoders

Today: NN, data
flow graph, and
data parallelism



Static Models vs. Dynamic Models



Static vs. Dynamic Dataflow Graphs

- Static Dataflow graphs
 - Define once, execute many times
 - Execution: Once defined, all following computation will **follow** the defined computation
 - Advantages
 - No extra effort for batching optimization, because it can be by nature batched
 - It is always easy to handle a static computational dataflow graphs in all aspects, because of its fixed structure
 - Node placement, distributed runtime, memory management, etc.
 - Benefit the developers

Static vs. Dynamic Dataflow Graphs

- Can we handle dynamic dataflow graphs?
 - Difficulty in expressing complex flow-control logic
 - Complexity of the computation graph implementation
 - Difficulty in debugging

How to Handle Dynamic Dataflow Graph?

- In general two ways:
 - Imperative: do not requiring contracting the entire graph before execution
 - Other symbolic representation on top of dataflow graph
 - vertex-centric representation

PYTORCH

DyNet


Chainer

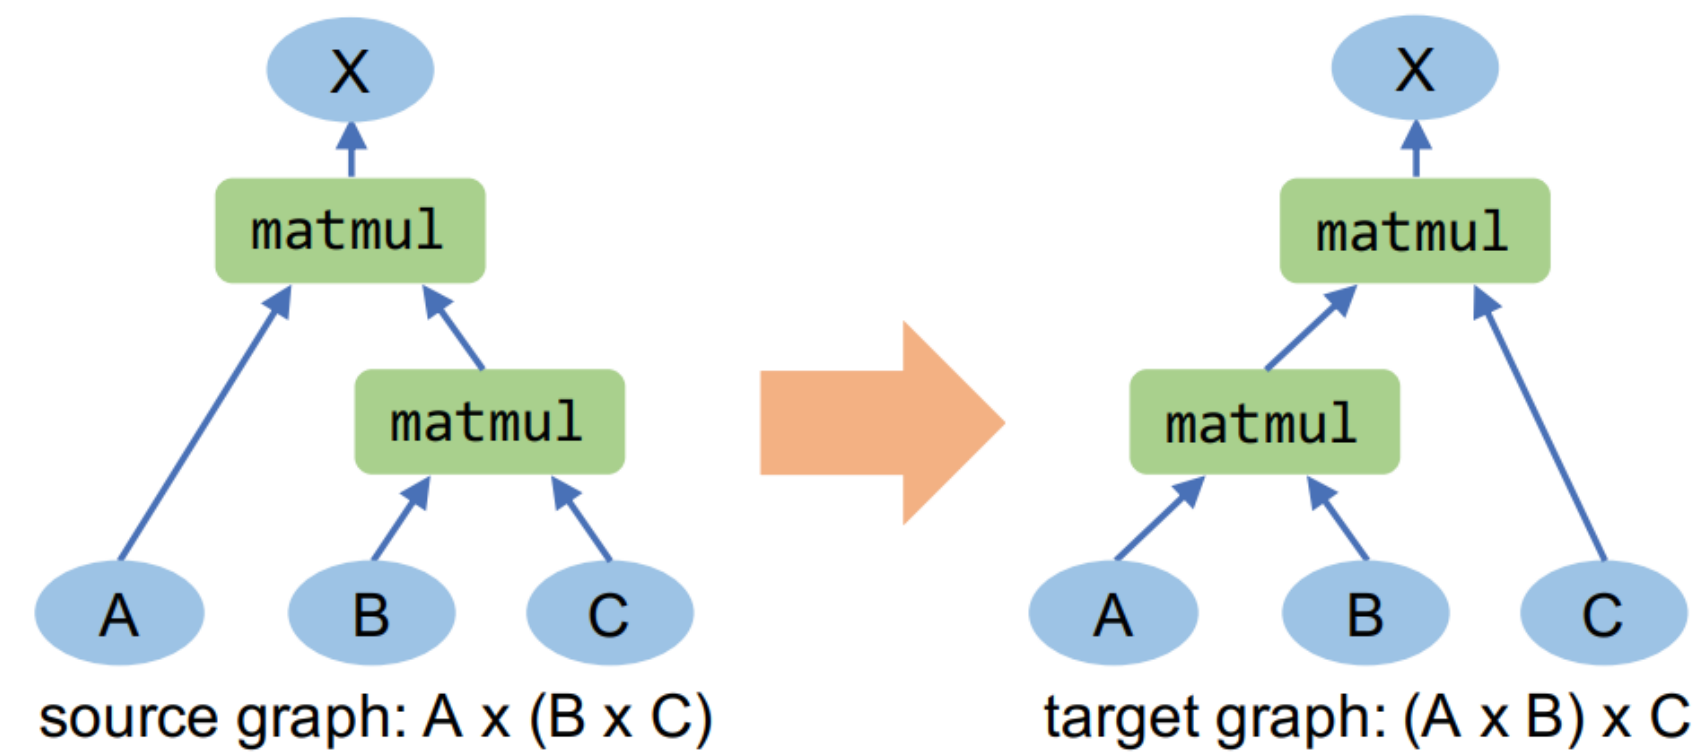
Imperative

Symbolic


Questions

- Is **CNN training** static or dynamic graph?
- Is **CNN inference** static or dynamic graph?
- Is **GPT-3 (transformers decoder)** training static graph or dynamic?
- Is **GPT-3 inference with batch size = 1** static or dynamic graph
- Is **GPT-3 serving** static or dynamic graph

Advanced Topic: DL Dataflow Graph Optimization



(a) Associativity of matrix multiplication.



(b) Fusing two matrix multiplications using concatenation and split.


Advanced Topic: DL Graph Compilation



High-level IR Optimizations and Transformations

Tensor Operator Level Optimization

↓ Direct code generation





Where We Are


- Deep Learning as Dataflow Graphs
- Auto-differentiation Libraries
 - Symbolic vs. Imperative
 - Static vs. Dynamic
- **DL Parallelism**

DL Parallelization: 3 Core Problems



Computing



Communication



Memory



$$\theta^{(t+1)} = f(\theta^{(t)}, \nabla_L(\theta^{(t)}, D^{(t)}))$$

parameter

weight update
(sgd, adam, etc.)

model
(CNN, GPT, etc.)

data

Two Views of ML Parallelisms

Classic view

Data parallelism

Model parallelism

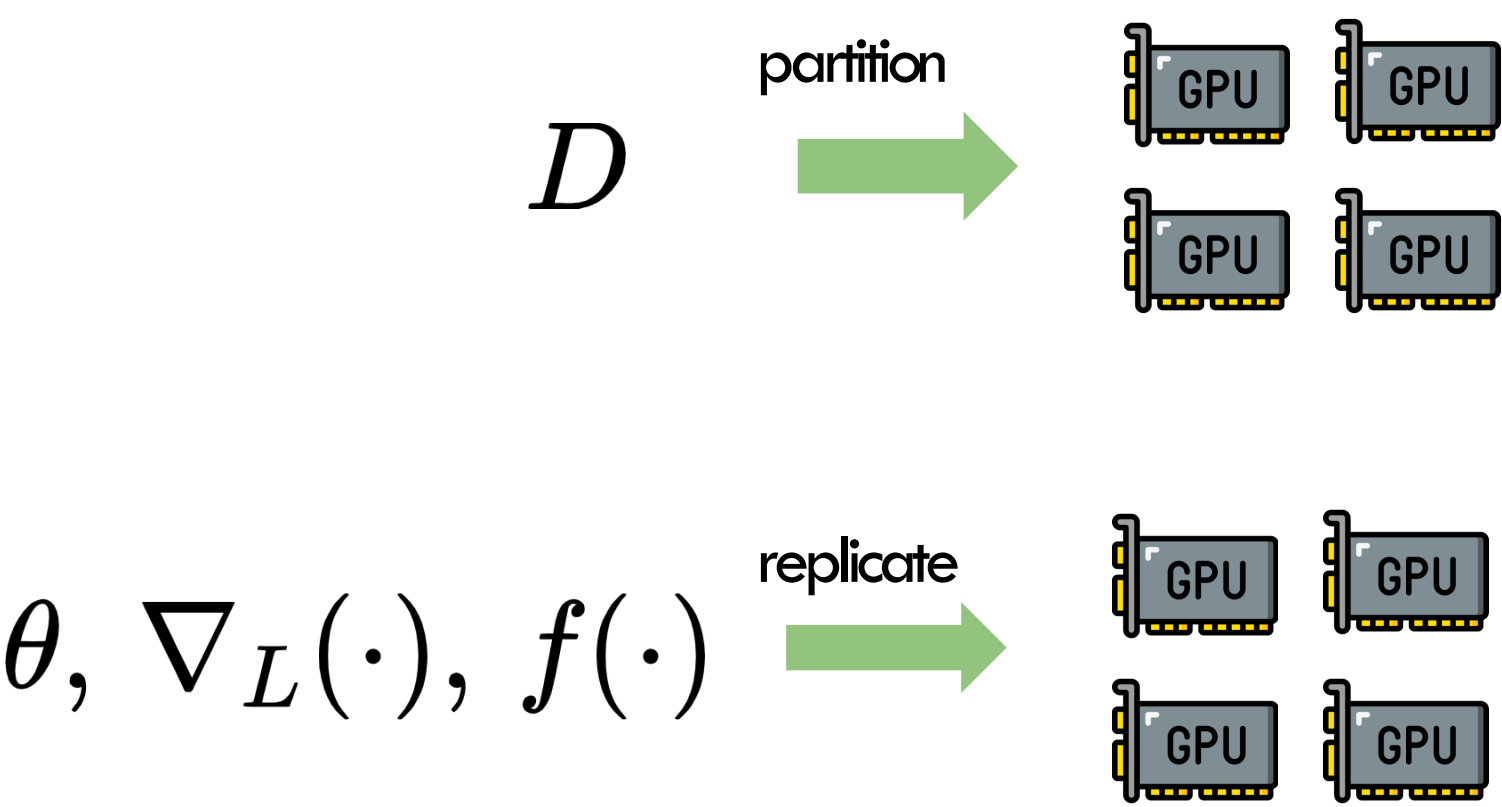
New view

Inter-op parallelism

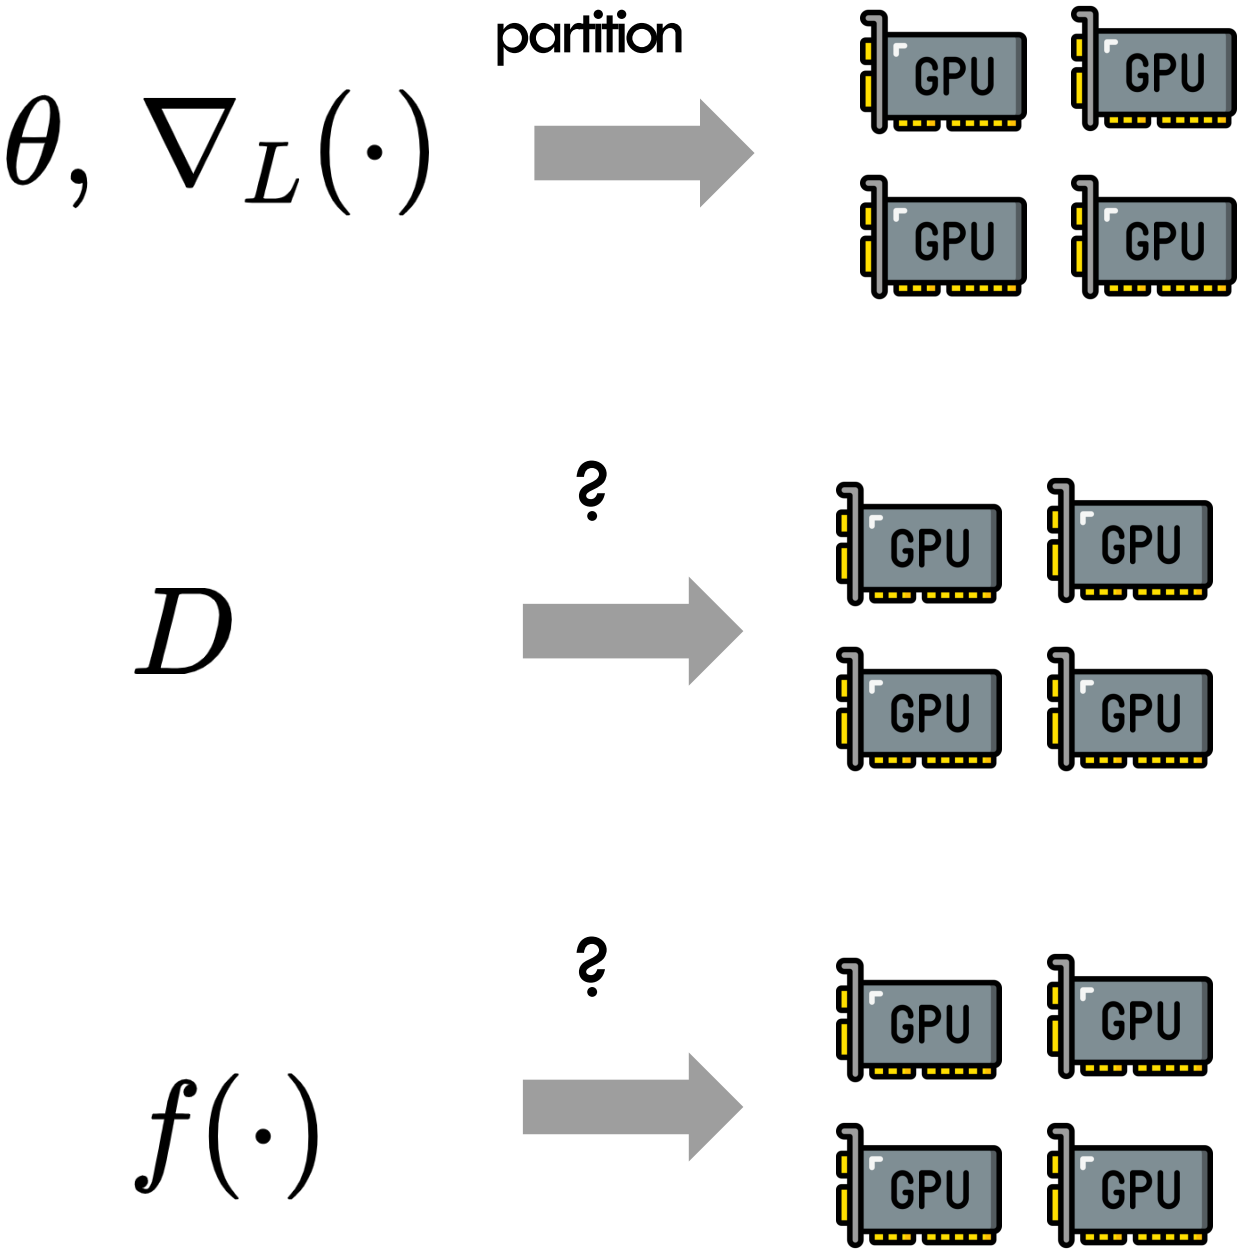
Intra-op parallelism

Data and Model Parallelism

Data parallelism



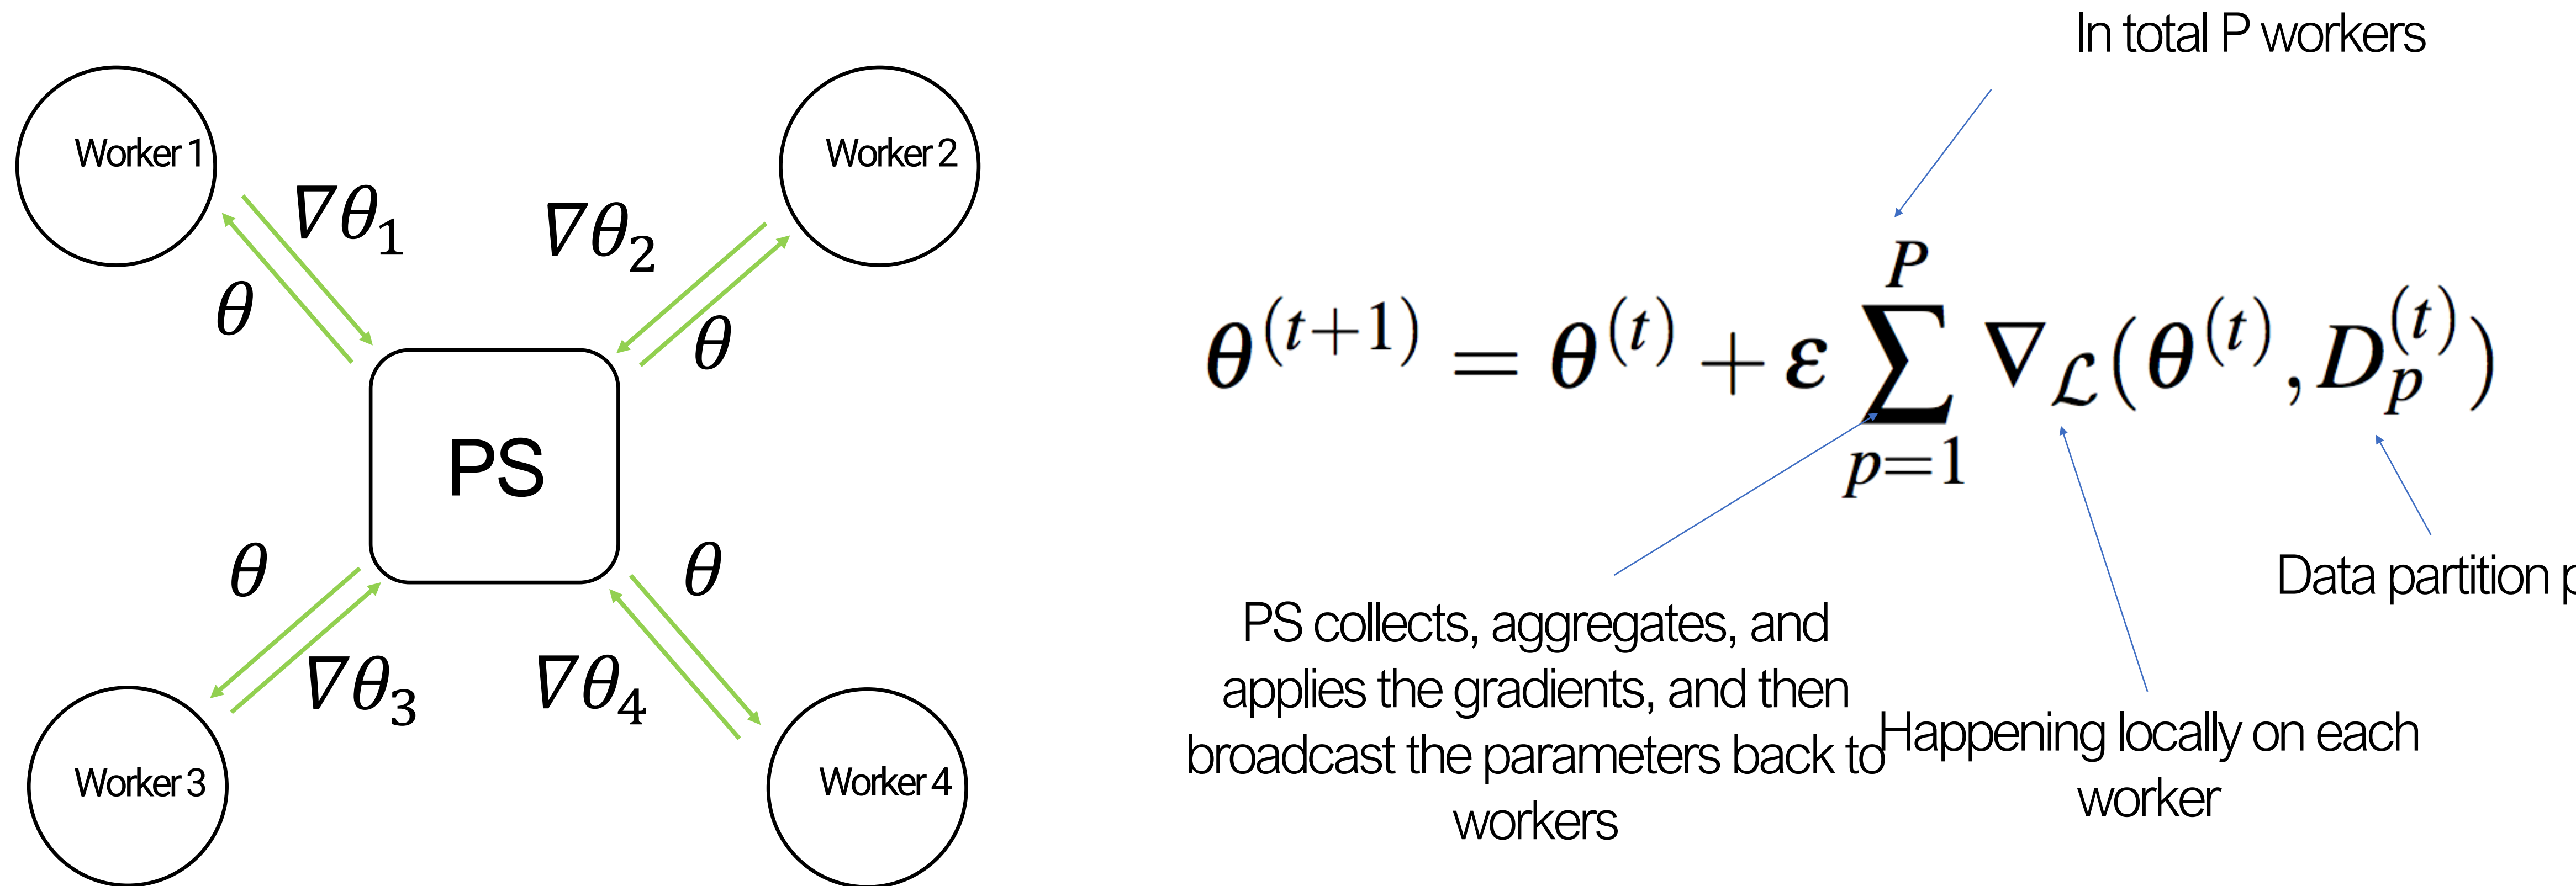
Model parallelism



$$\theta^{(t+1)} = f(\theta^{(t)}, \nabla_L(\theta^{(t)}, D^{(t)}))$$

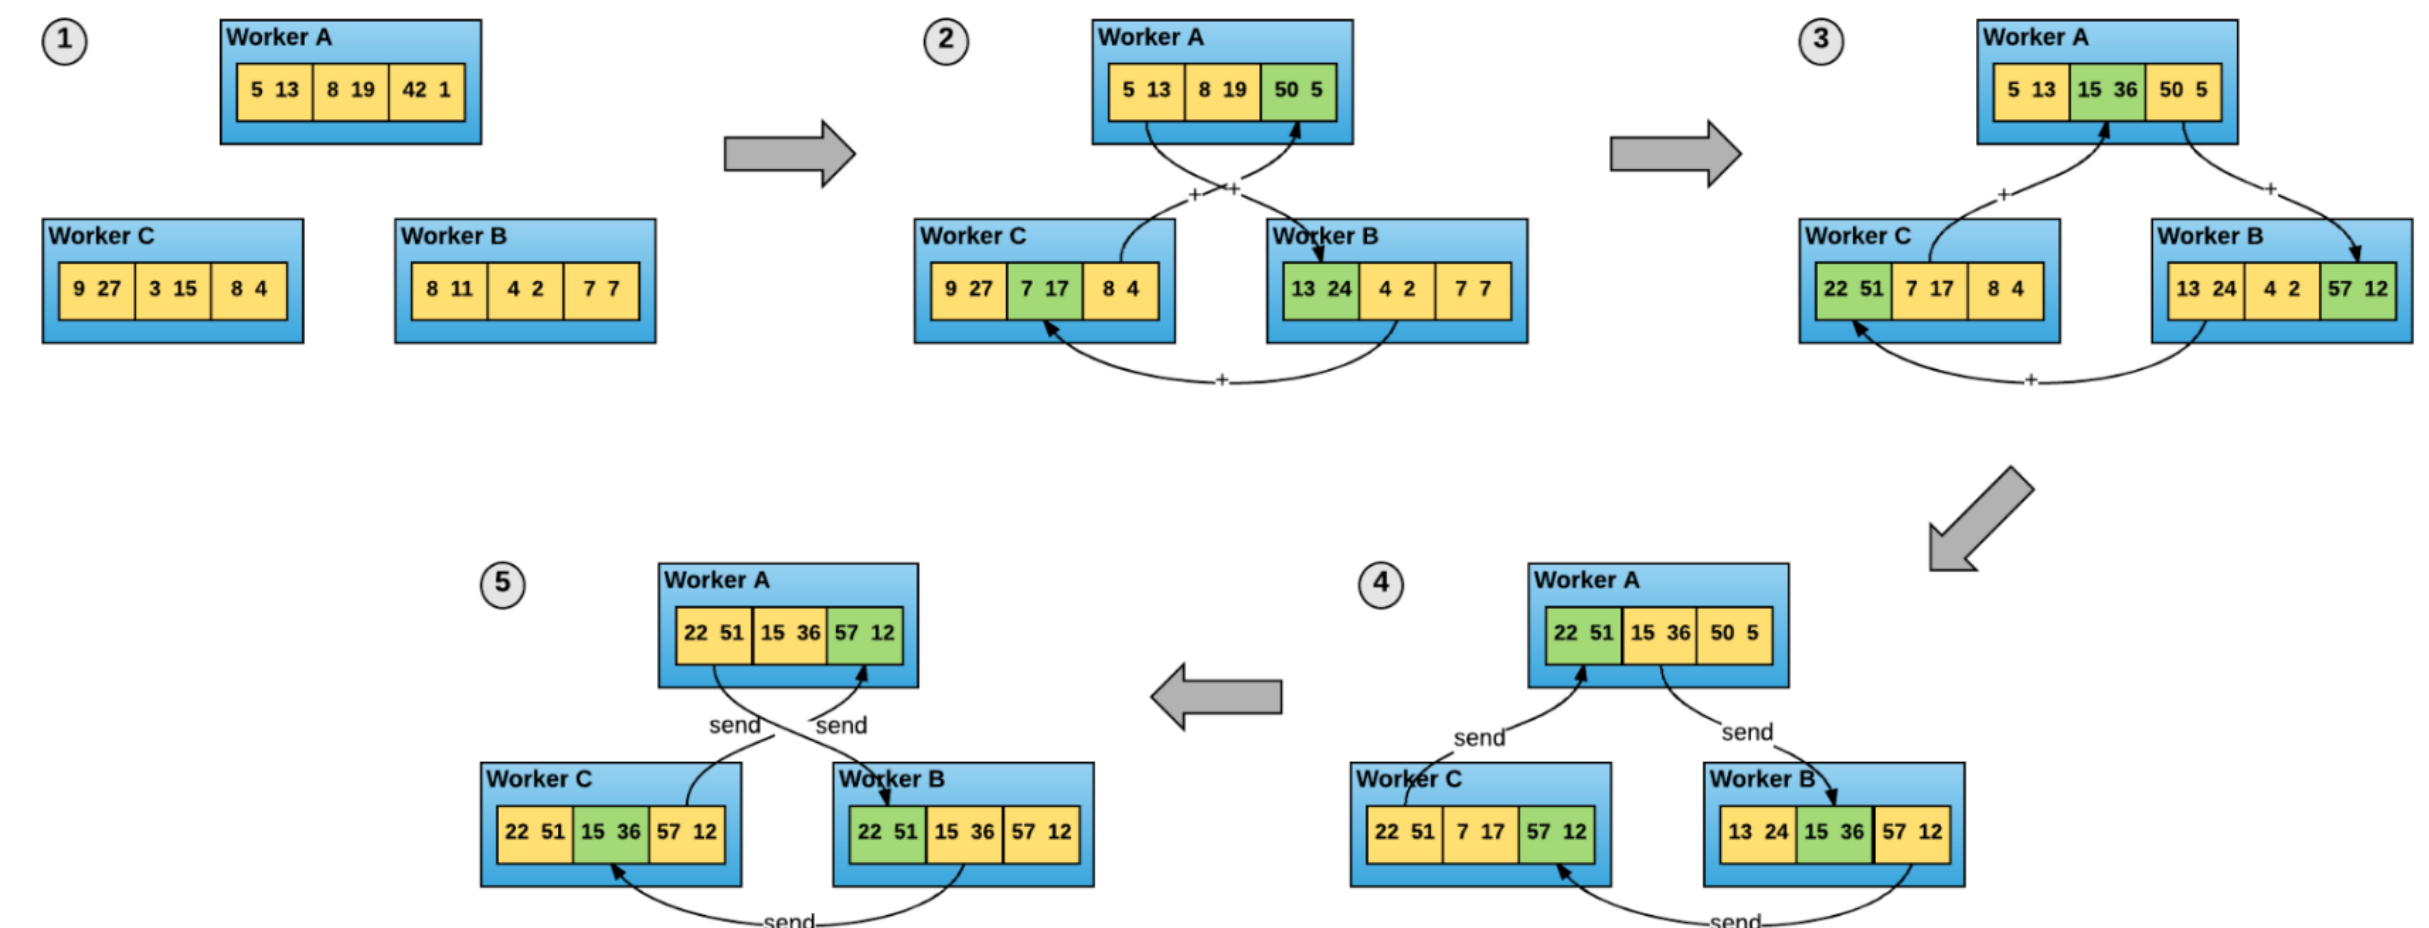

parameter weight update (sgd, adam, etc.) model (CNN, GPT, etc.) data

PS Implements Data Parallelism




Representee Systems: Poseidon, GeePS, BytePS, etc.

AllReduce Can Also Handle Data Parallelism Comm



Representee Systems: Horovod, Torch.DDP

Big Model: The Core Computational Challenge






How to train and serve big models?



Model Parallelism

Two Views of ML Parallelisms

Data and model parallelism

- Two pillars: **data** and **model**.
-  “Data parallelism” is general and precise.
-  “Model parallelism” is vague.
-  The view creates ambiguity for methods that neither partitions data nor the model computation.

New: Inter-op and Intra-op parallelism.

- Two pillars: **computational graph** and **device cluster**
-  This view is based on their computing characteristics.
-  This view facilitates the development of new parallelism methods.

Device Cluster

Nvidia DGX with V100

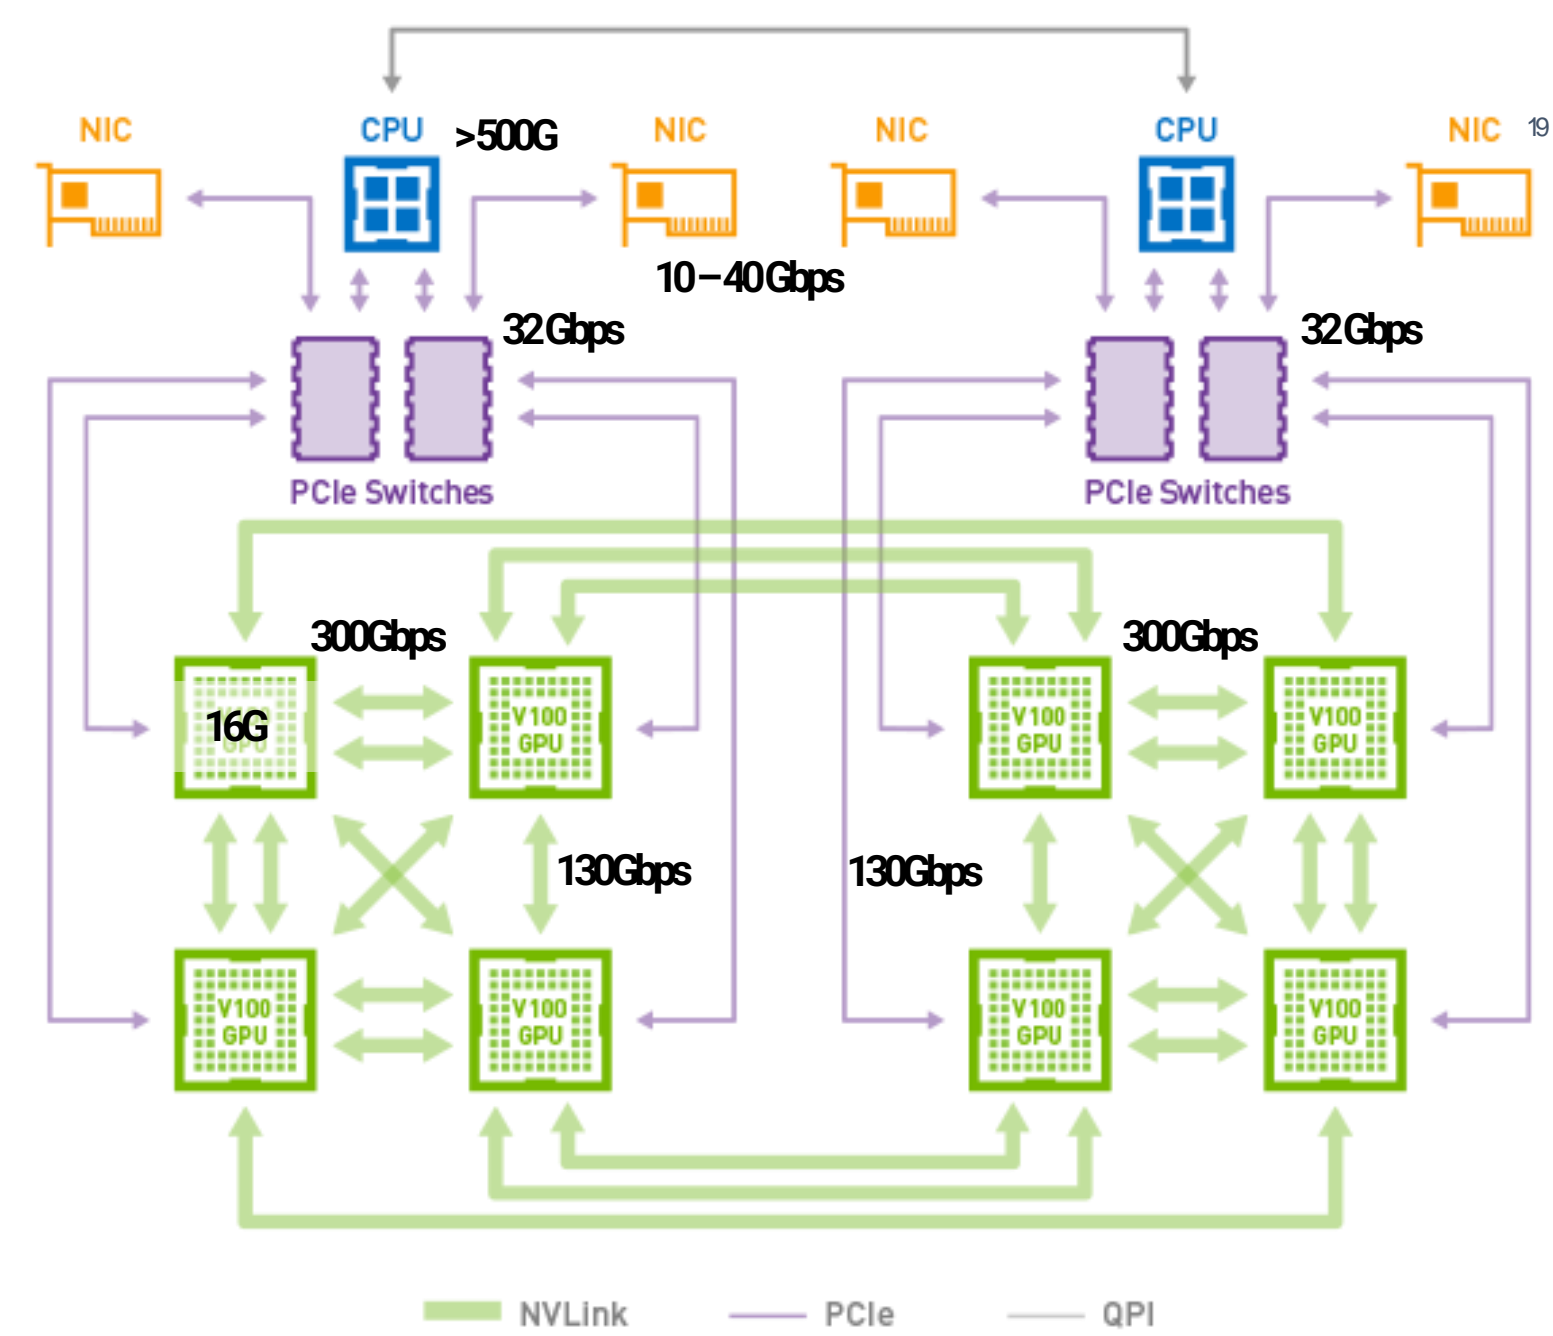




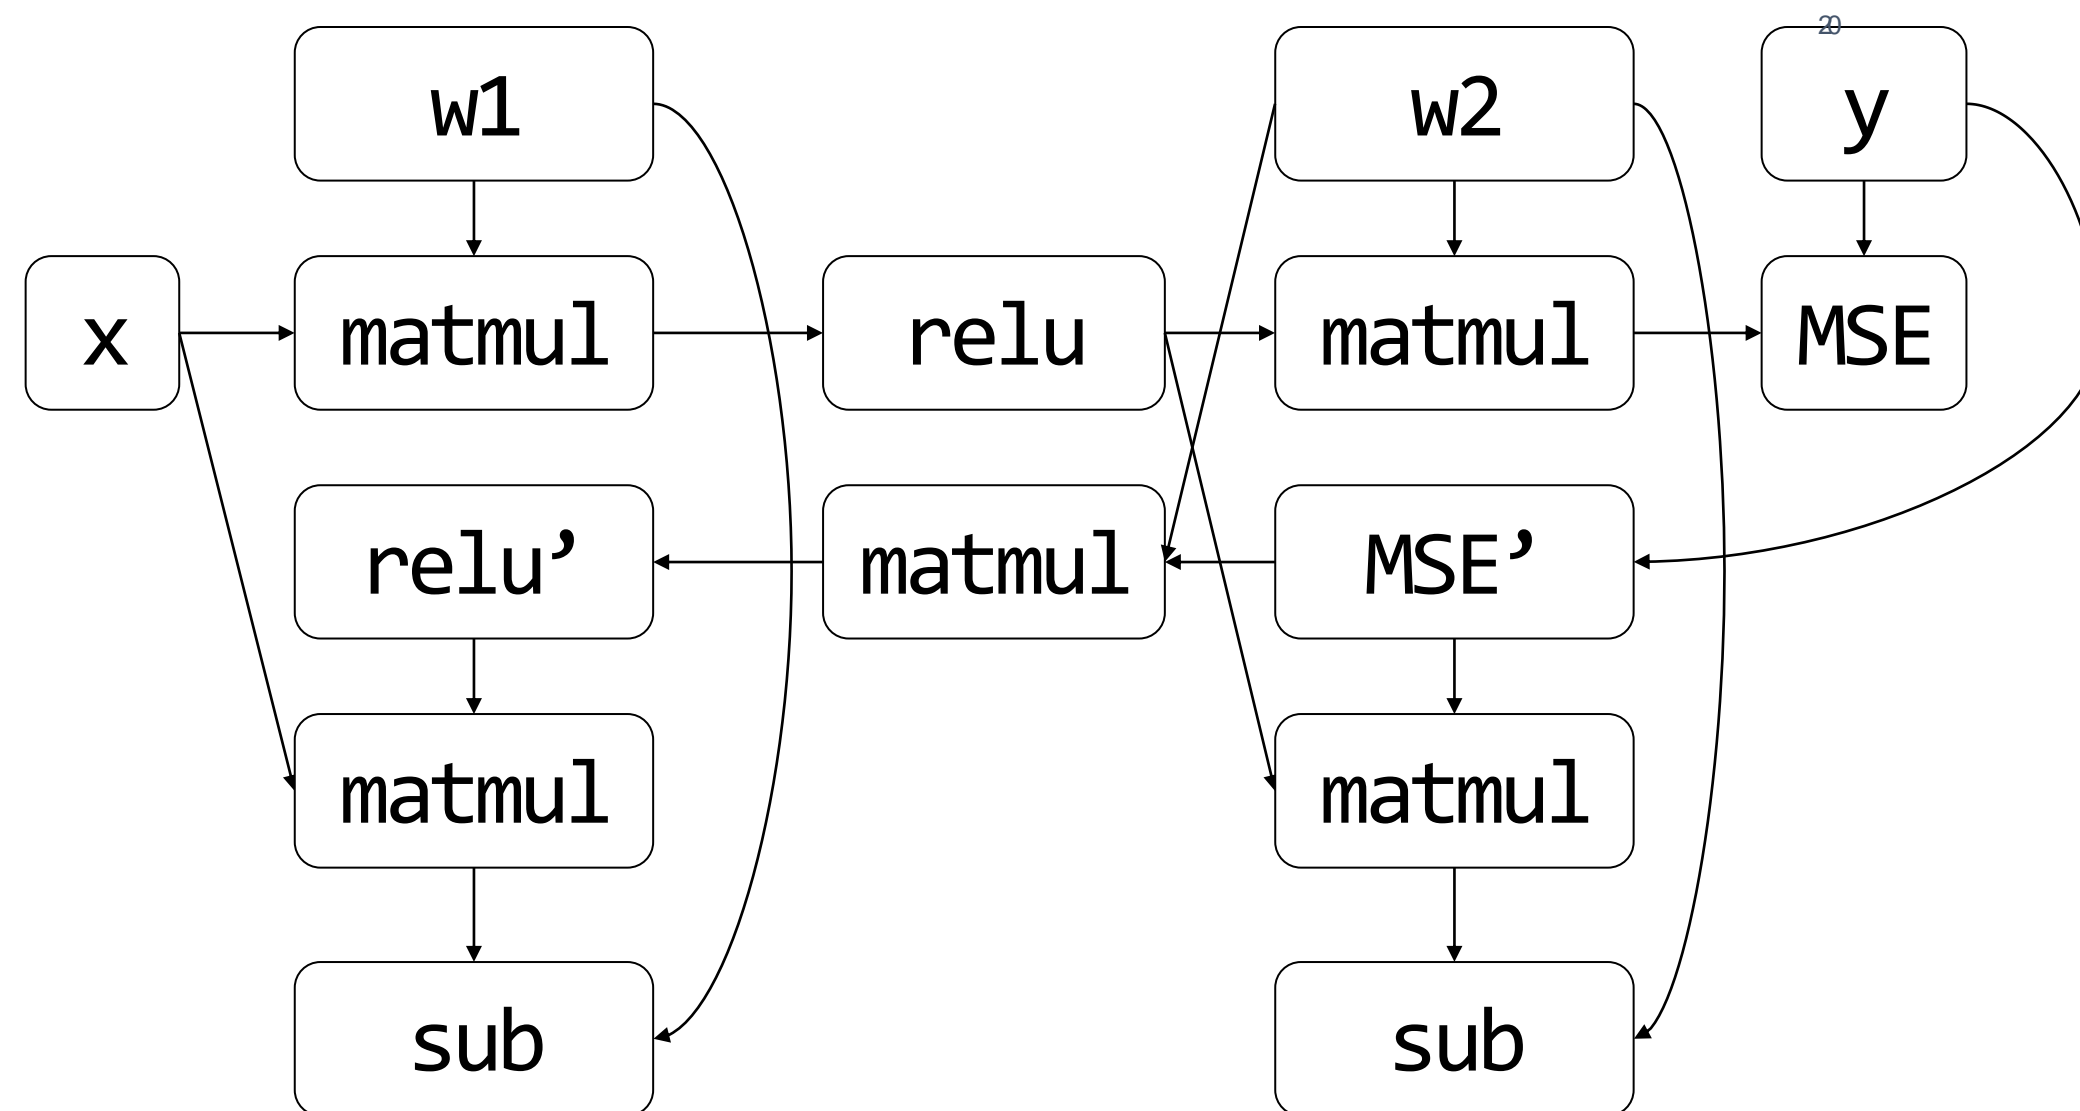
Figure from NVIDIA

A typical GPU cluster topology

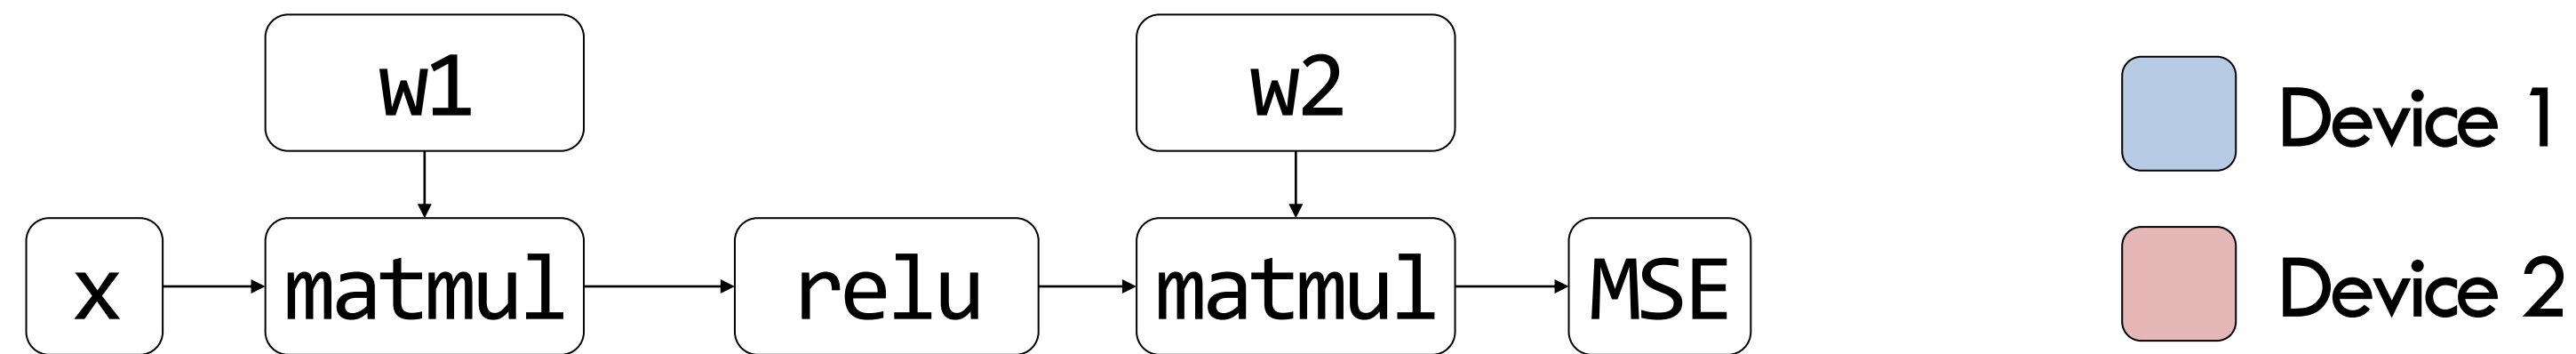


Partitioning Computation Graph on Device Cluster



How to partition the computational graph on the device cluster?





Partitioning Computation Graph





Strategy 1





Strategy 2





Partitioning Computation Graph



Strategy 1: Inter-operator Parallelism




Strategy 2: Intra-operator Parallelism


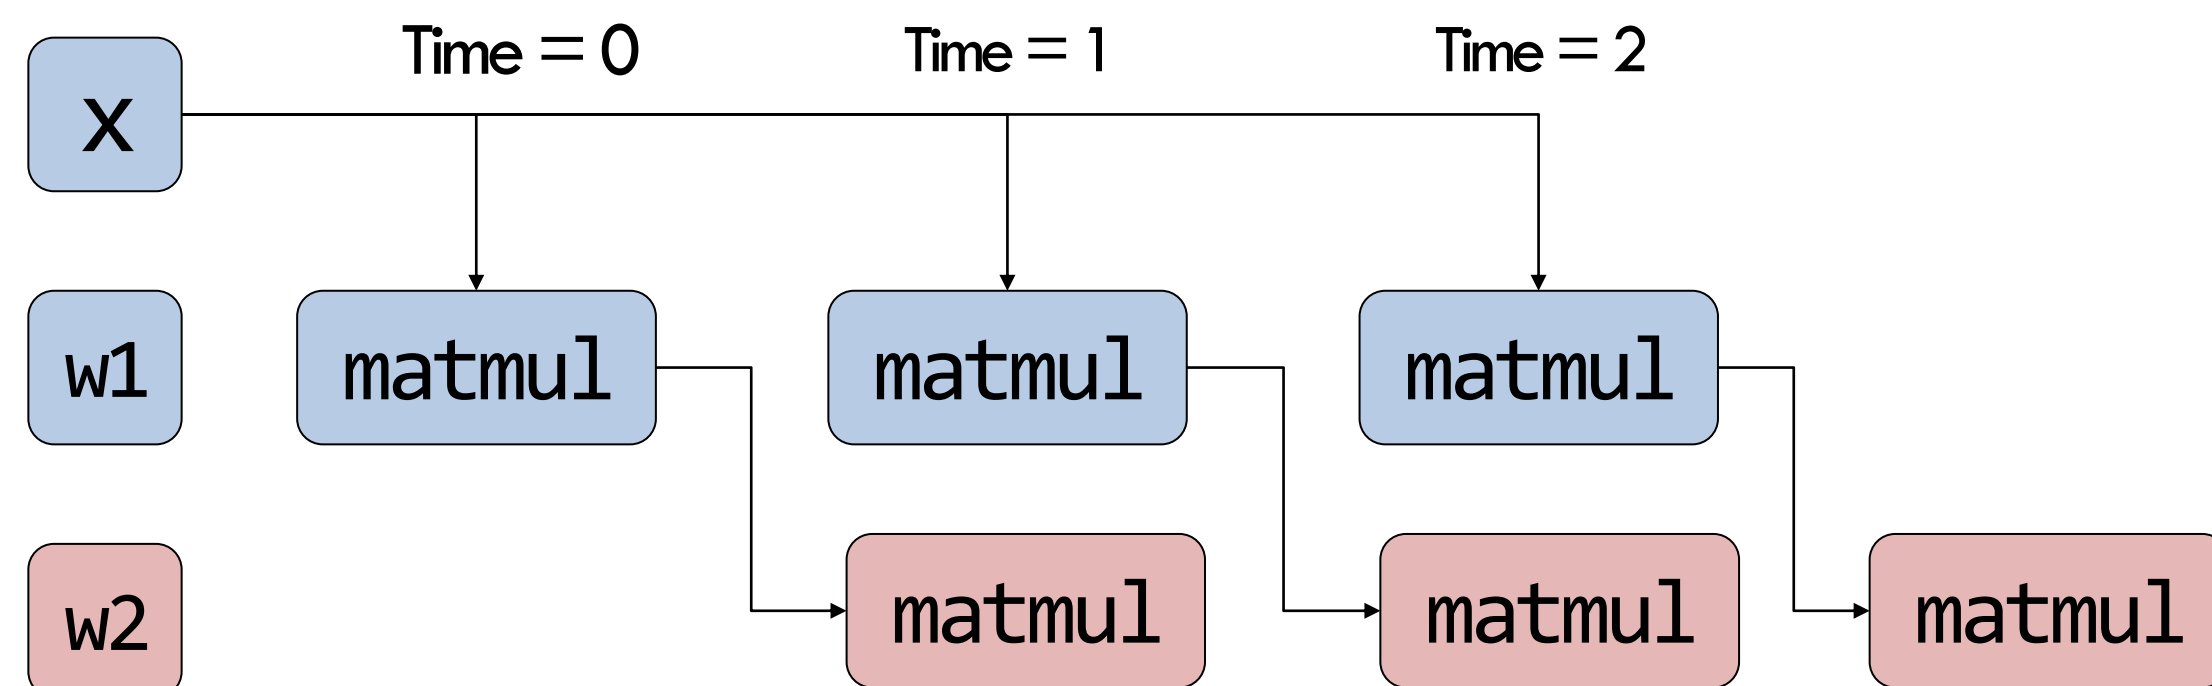




More Parallelisms...


Multiple intra-op strategies for a single node




More strategies




Summary: Inter-op and Intra-op Parallelisms



Inter-op parallelism: Assign different operators to different devices.






Intra-op parallelism: Assign different regions of a single operator to different devices.



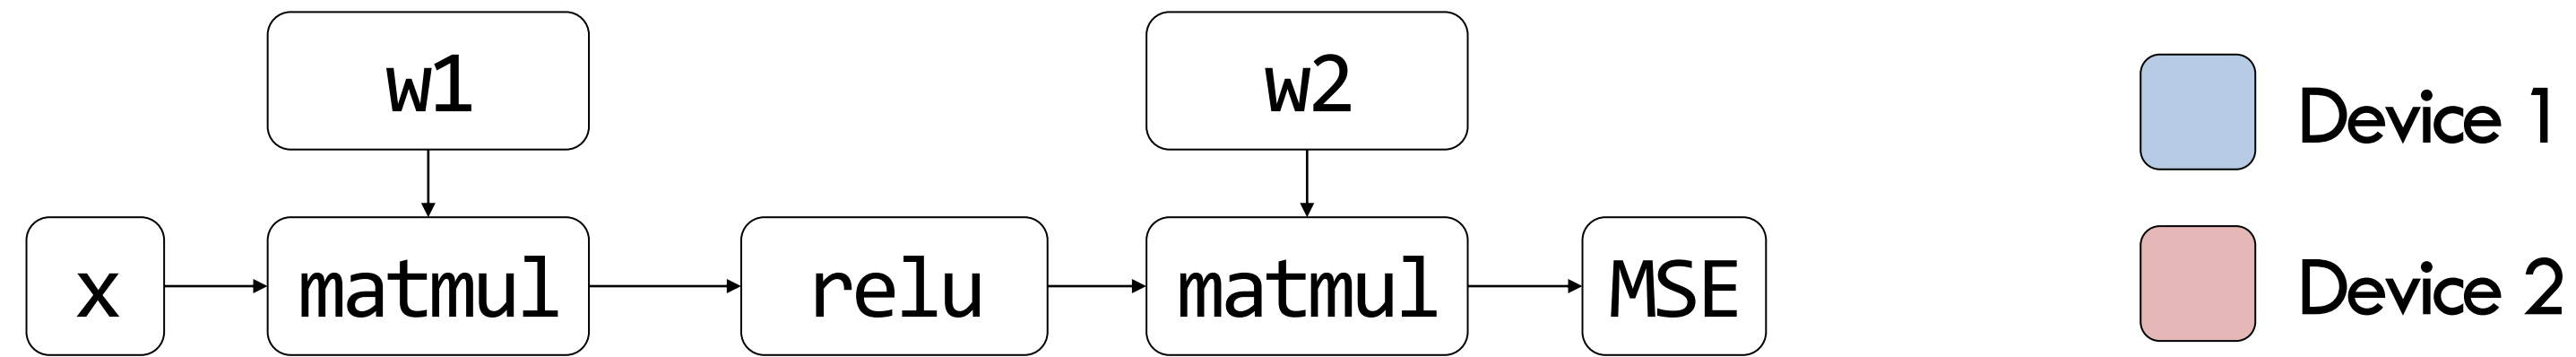
Inside Intra- and Inter-op Parallelism



$$Y = X \cdot W_1 \cdot W_2 = X \cdot \begin{bmatrix} W_1^{d1} & W_1^{d2} \end{bmatrix} \cdot \begin{bmatrix} W_2^{d1} \\ W_2^{d2} \end{bmatrix}$$




Inter-op and Intra-op Parallelism: Characteristics



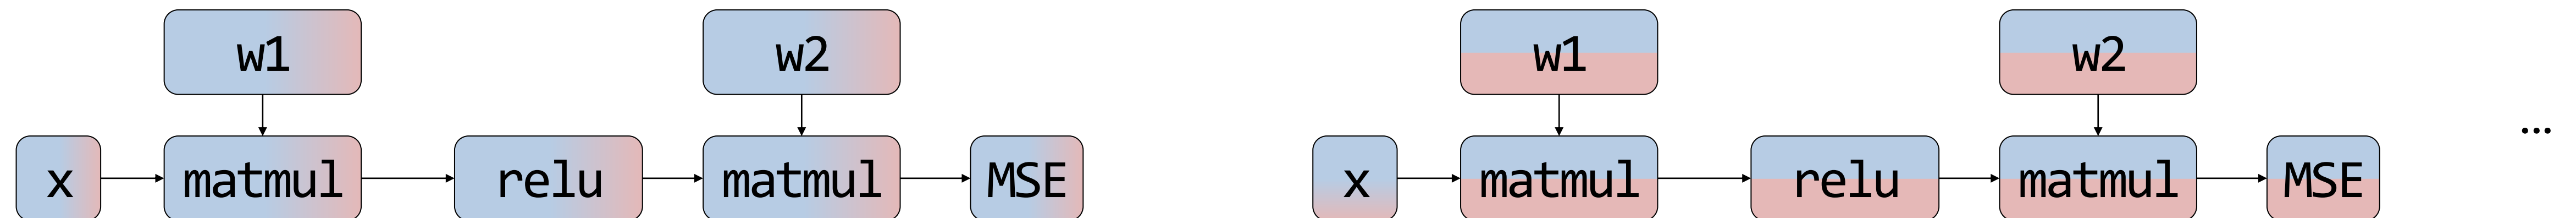
Inter-op parallelism:

Requires point-to-point communication but results in device idle




Intra-op parallelism:

Devices are busy but requires collective communication




Inter-op and Intra-op Parallelism: Characteristics



Device 1


Device 2

Inter-op parallelism



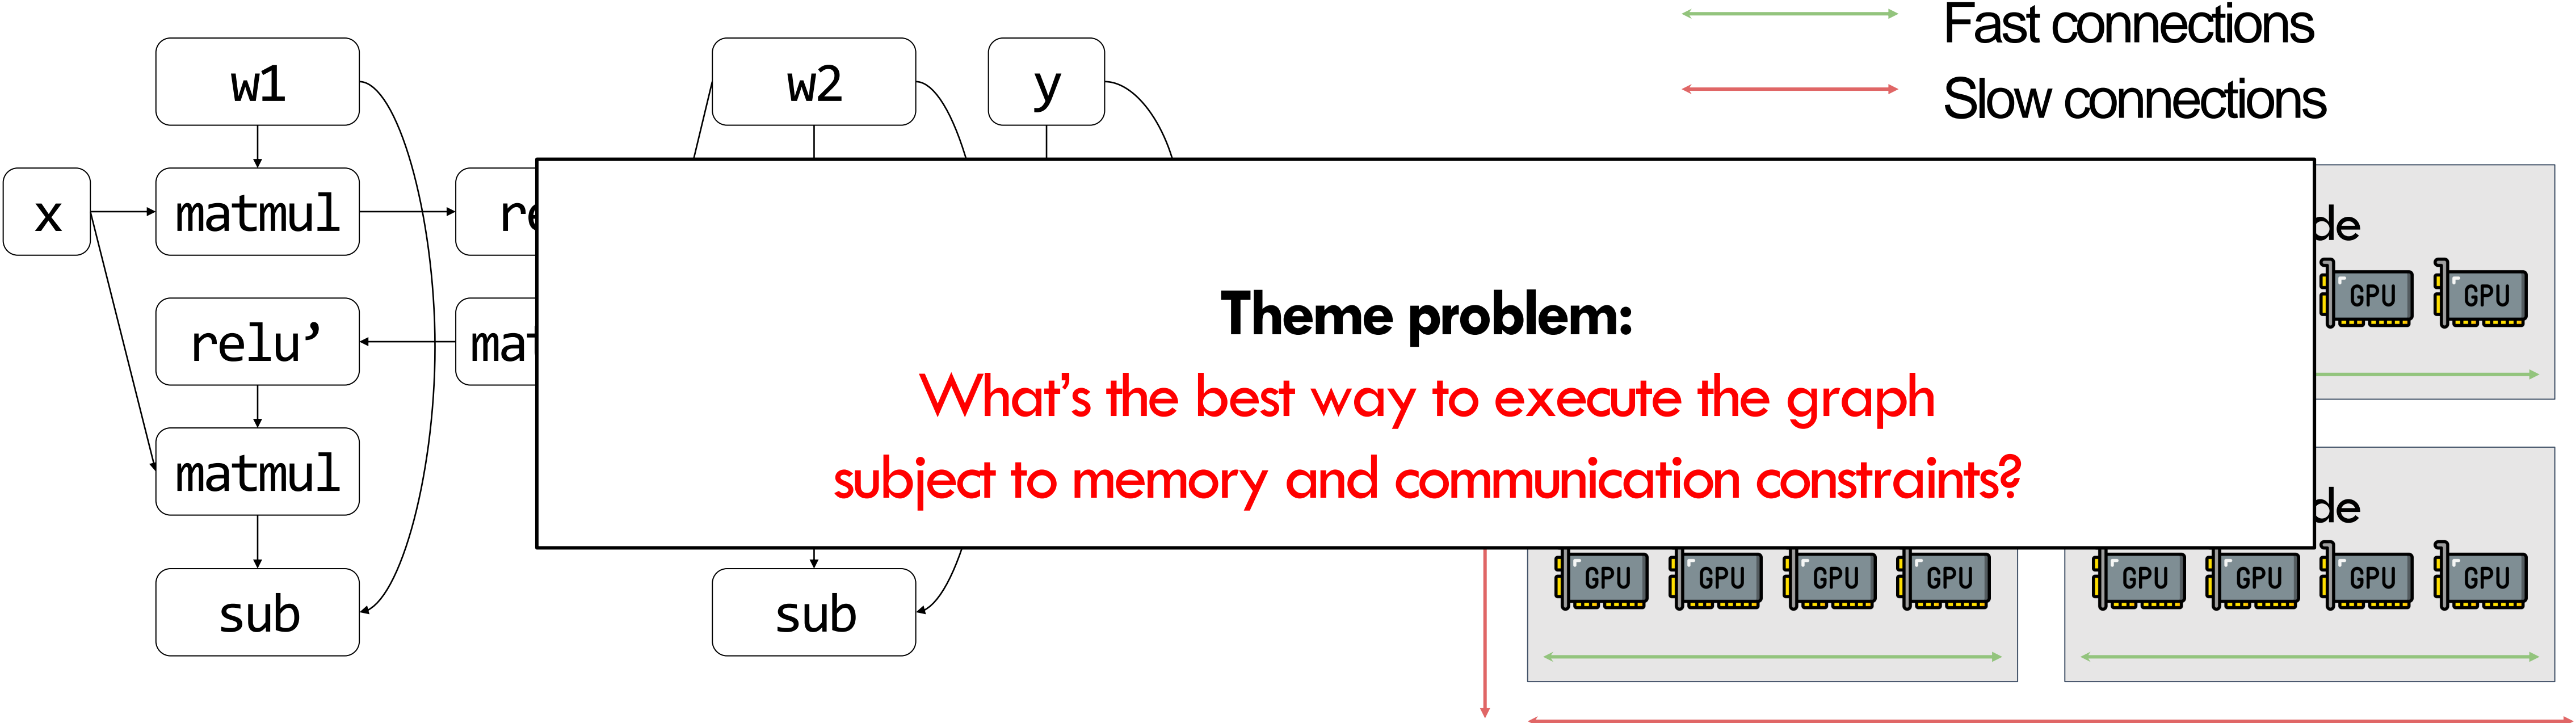
Trade-off

Intra-op parallelism



	Inter-operator Parallelism	Intra-operator Parallelism
Communication	Less	More
Device Idle Time	More	Less

ML Parallelization under New View




Where We Are


- Deep Learning as Dataflow Graphs
- Auto-differentiation Libraries
 - Symbolic vs. Imperative
 - Static vs. Dynamic
- DL Parallelism
 - **Inter-op parallelism**
 - Intra-op parallelism

Computational Graph (Neural Networks) \rightarrow Stages

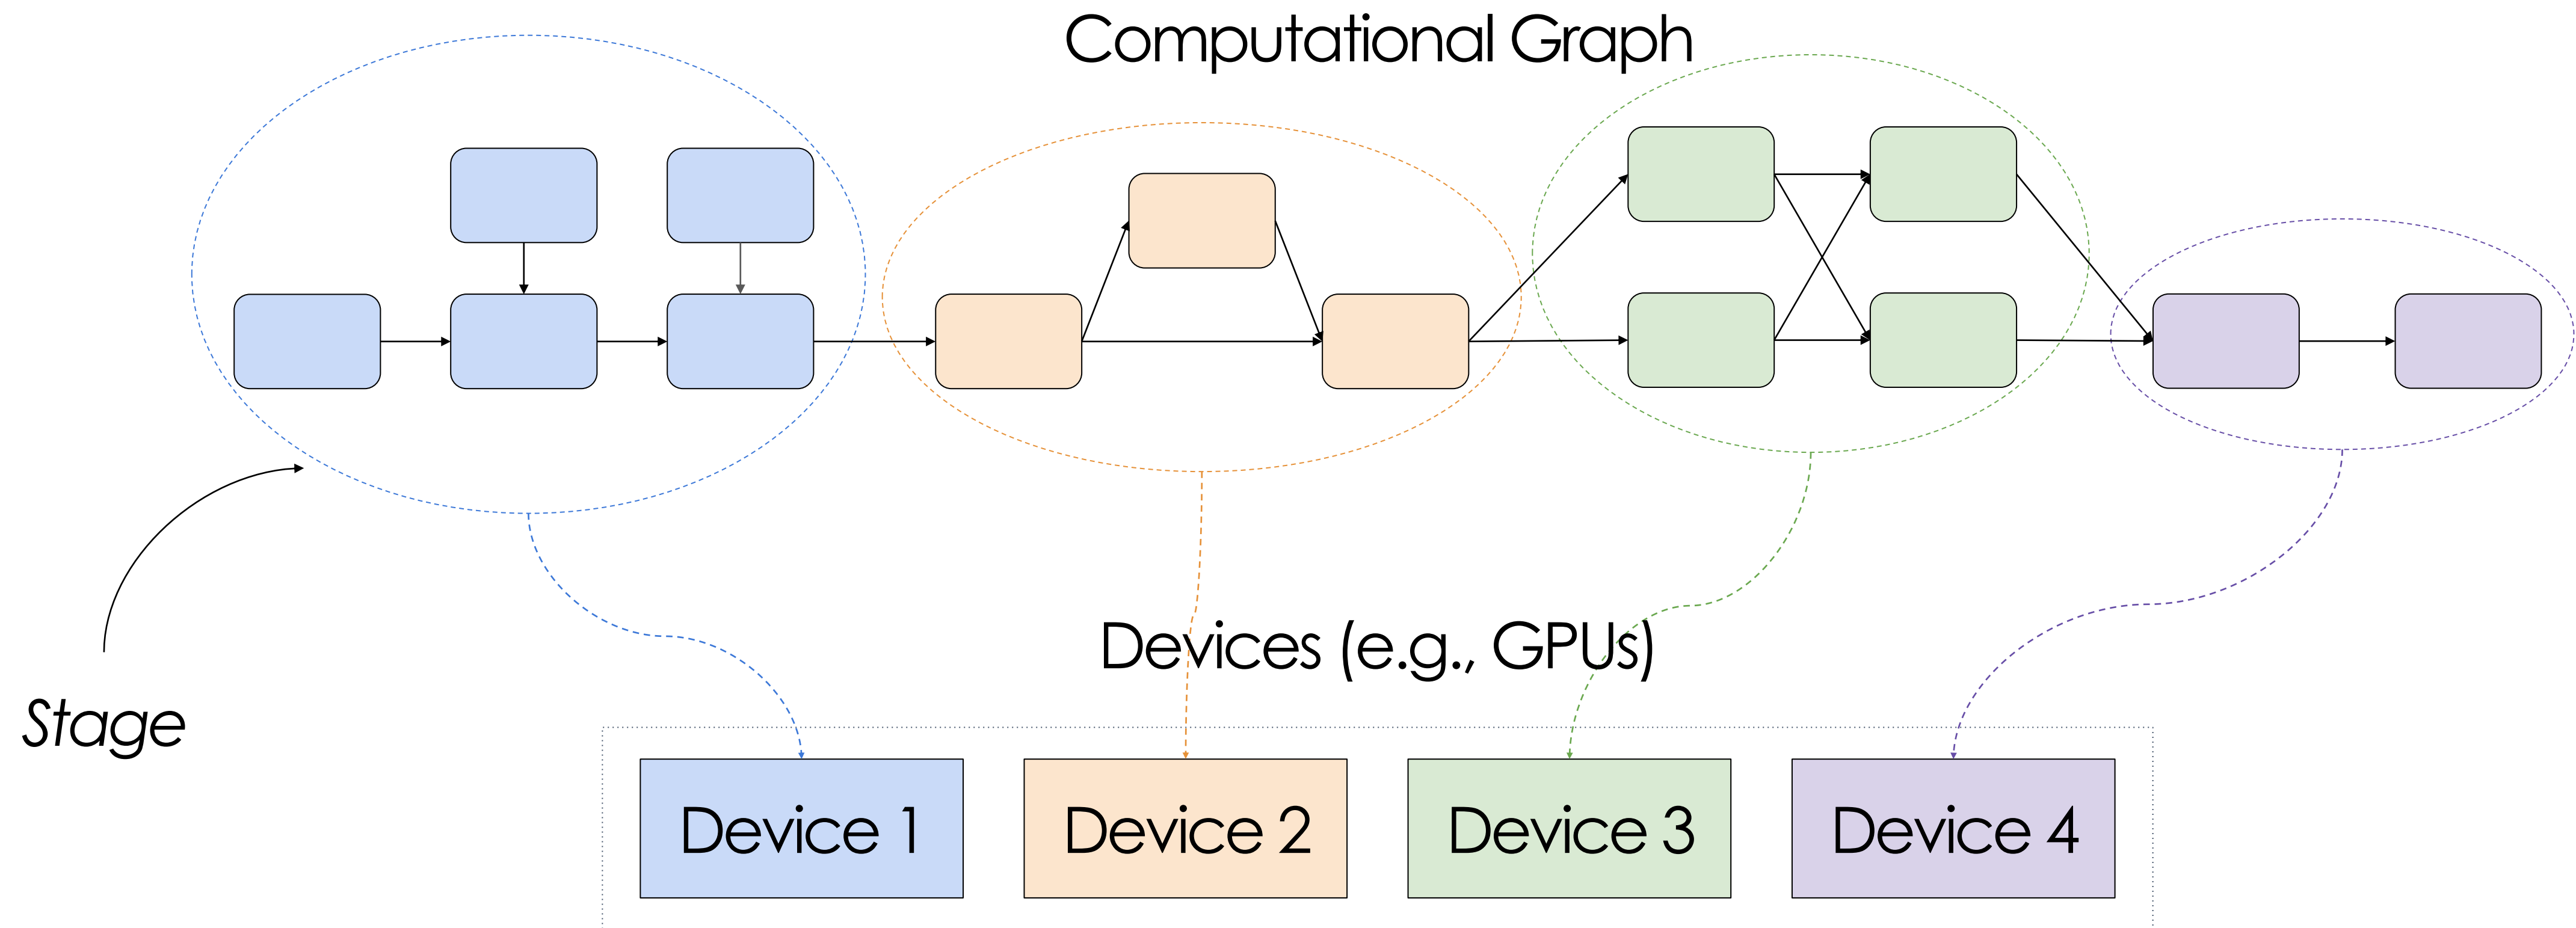
Computational Graph



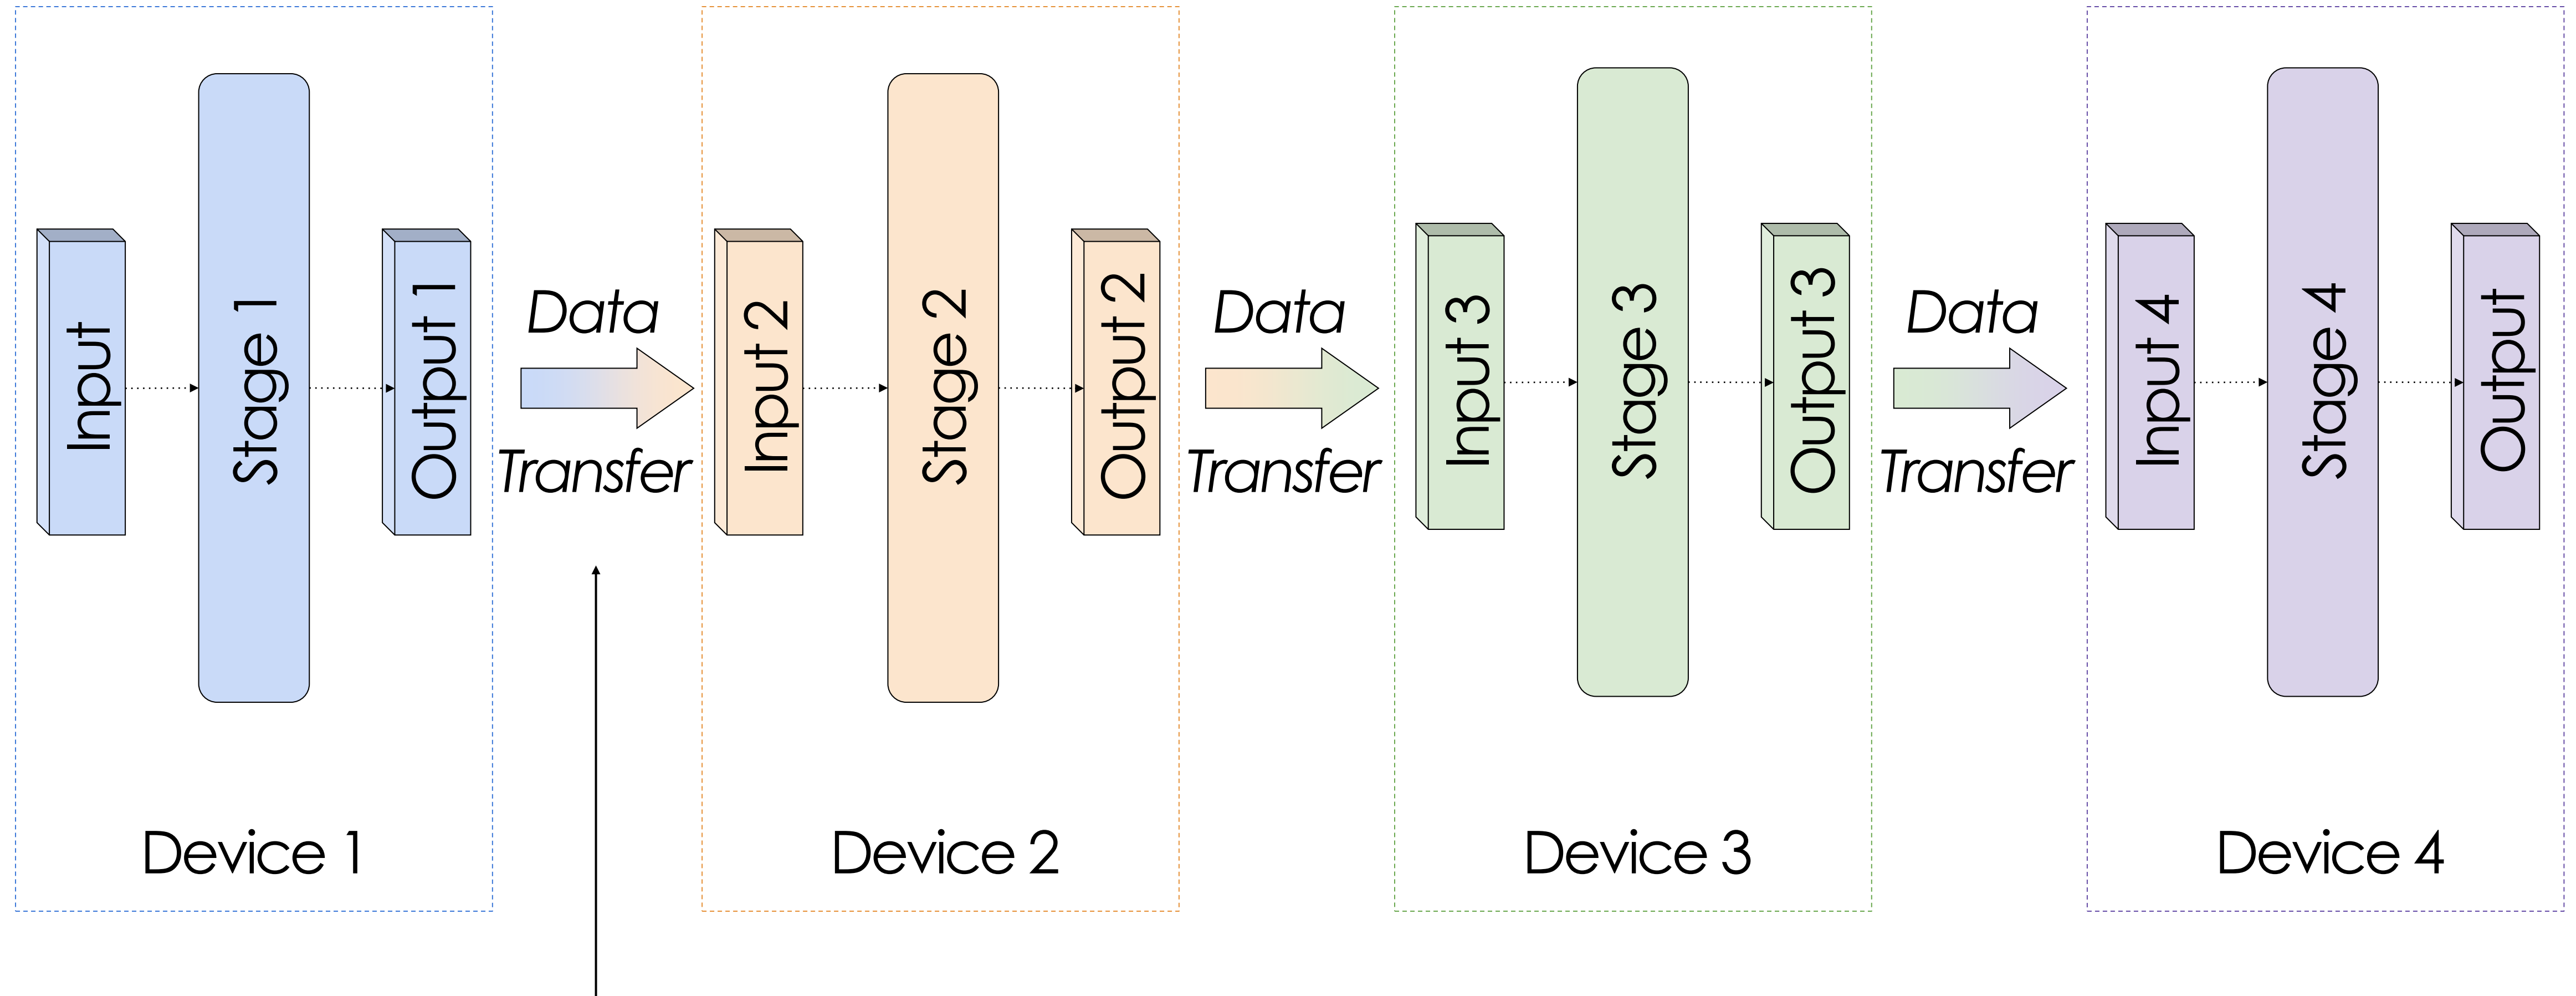
Devices (e.g., GPUs)



Computational Graph (Neural Networks) → Stages





Execution & Data Movement




Note: The time spent on data transfer is typically **small**, since we only communicate stage outputs at stage boundaries between two stages.

Timeline: Visualization of Inter-Operator Parallelism

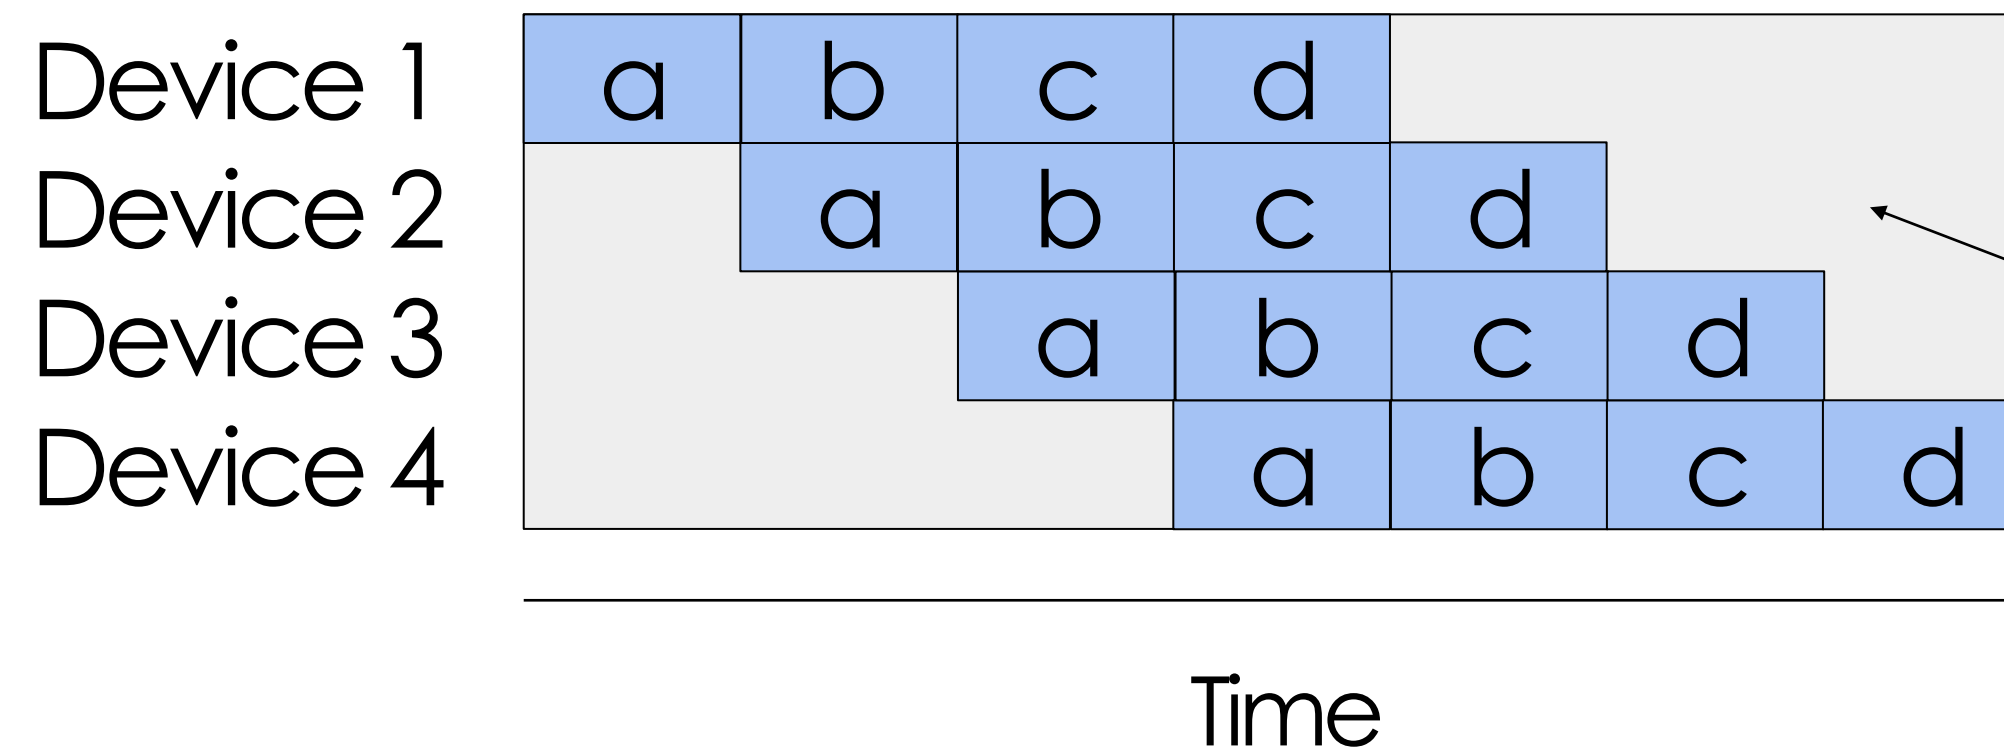


- Gray area () indicates devices being idle (a.k.a. Pipeline bubbles).
- Only 1 device activated at a time.
- **Pipeline bubble percentage** = $\text{bubble_area} / \text{total_area}$
= $(D - 1) / D$, assuming D devices.

Reduce Pipeline Bubbles via Pipelining Inputs

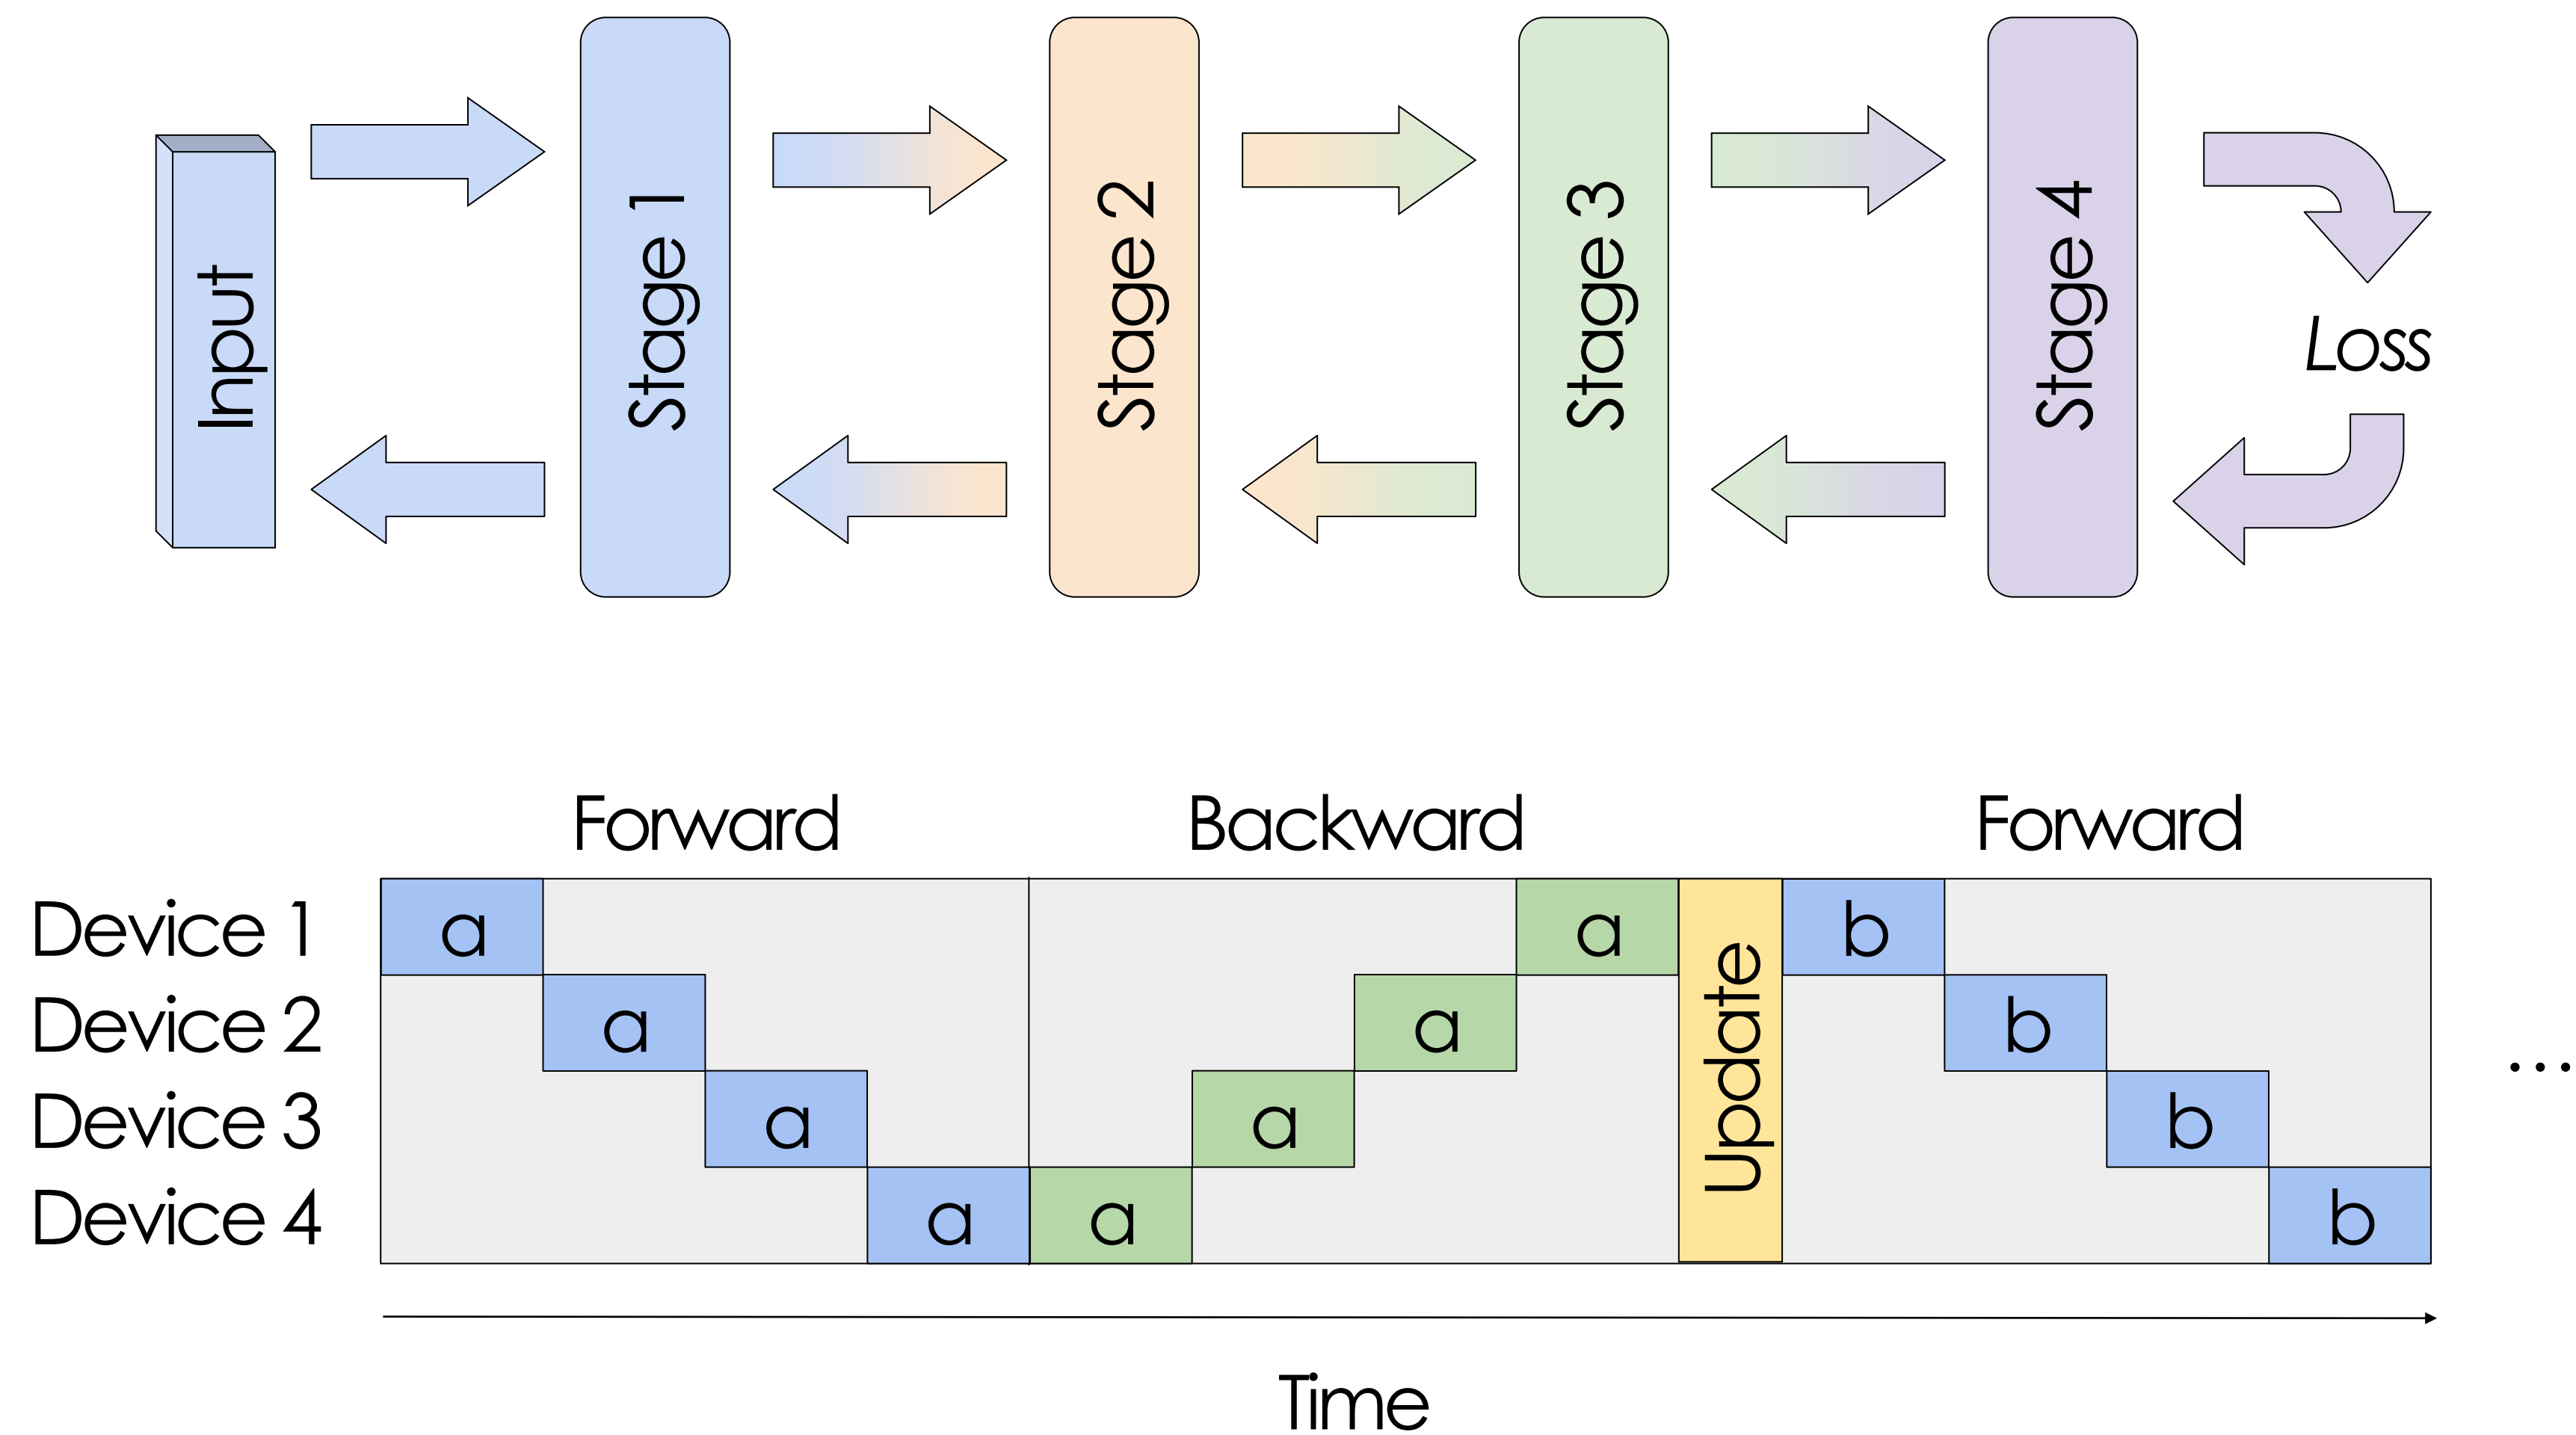


Used in inference.



Pipeline bubbles percentage
 $= (D - 1) / (D - 1 + N)$
with D devices and N inputs.

Training: Forward & Backward Dependency



How to Reduce Pipeline Bubbles for Training?

- Synchronous Pipeline Parallel Algorithms
 - GPipe
 - 1F1B
 - Interleaved 1F1B
 - TeraPipe
 - Chimera
- Asynchronous Pipeline Parallel Algorithms
 - AMPNet
 - Pipedream/Pipedream-2BW

How to Reduce Pipeline Bubbles for Training?

- Synchronous Pipeline Parallel Algorithms
 - **GPipe**
 - **1F1B**
 - Interleaved 1F1B
 - TeraPipe
 - Chimera
- Asynchronous Pipeline Parallel Algorithms
 - AMPNet
 - Pipedream/Pipedream-2BW

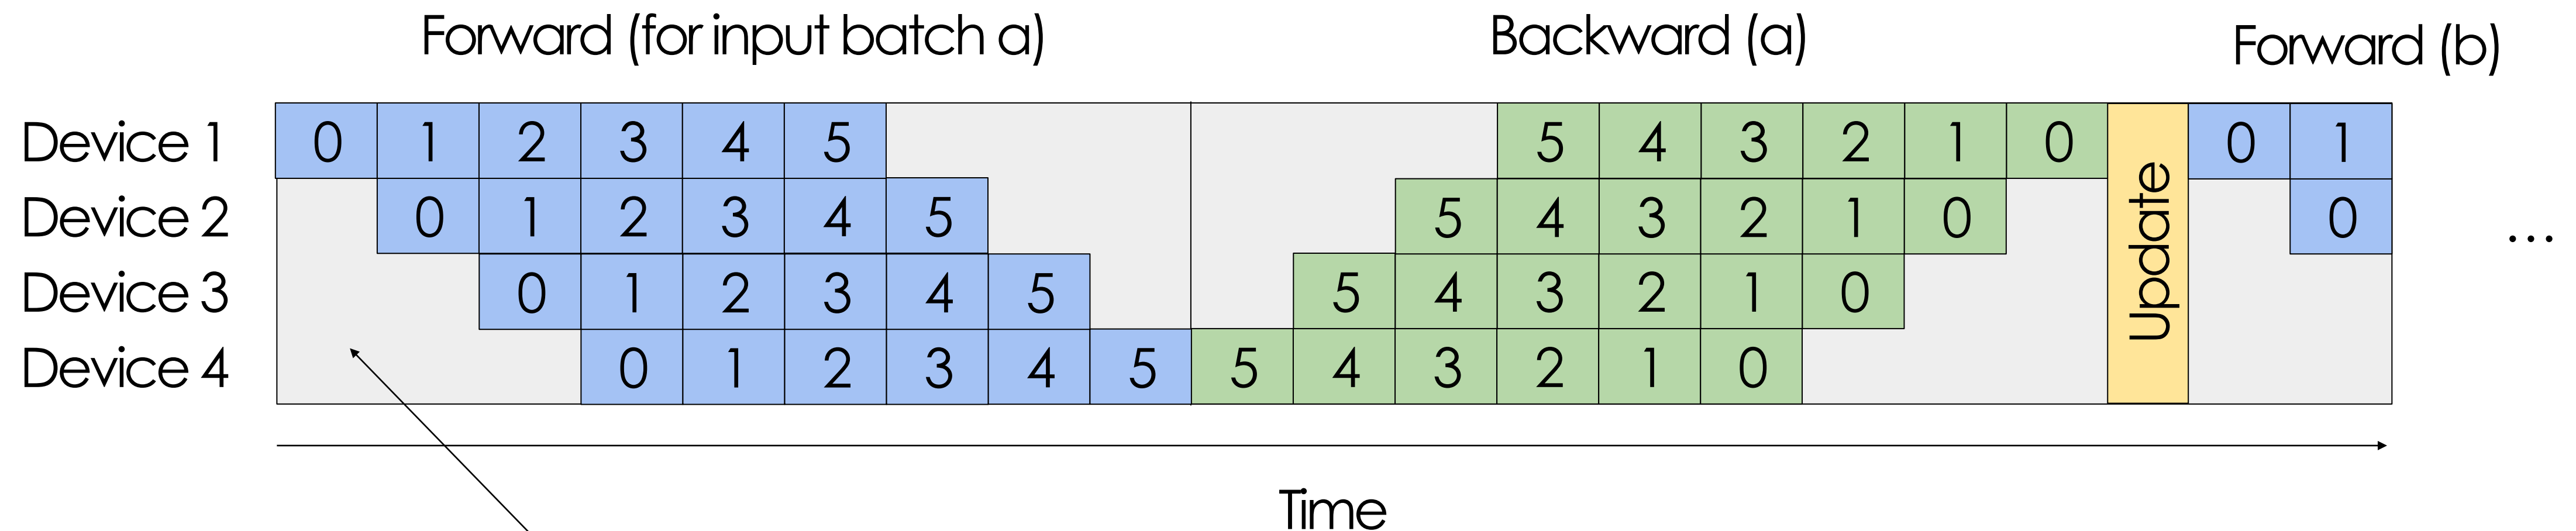
GPipe

Idea: Partition the input batch into multiple “*micro-batches*”. Pipeline the micro-batches.

Accumulate the gradients of the micro-batches:

$$\nabla L_{\theta}(x) = \frac{1}{N} \sum_{i=1}^N \nabla L_{\theta}(x_i)$$

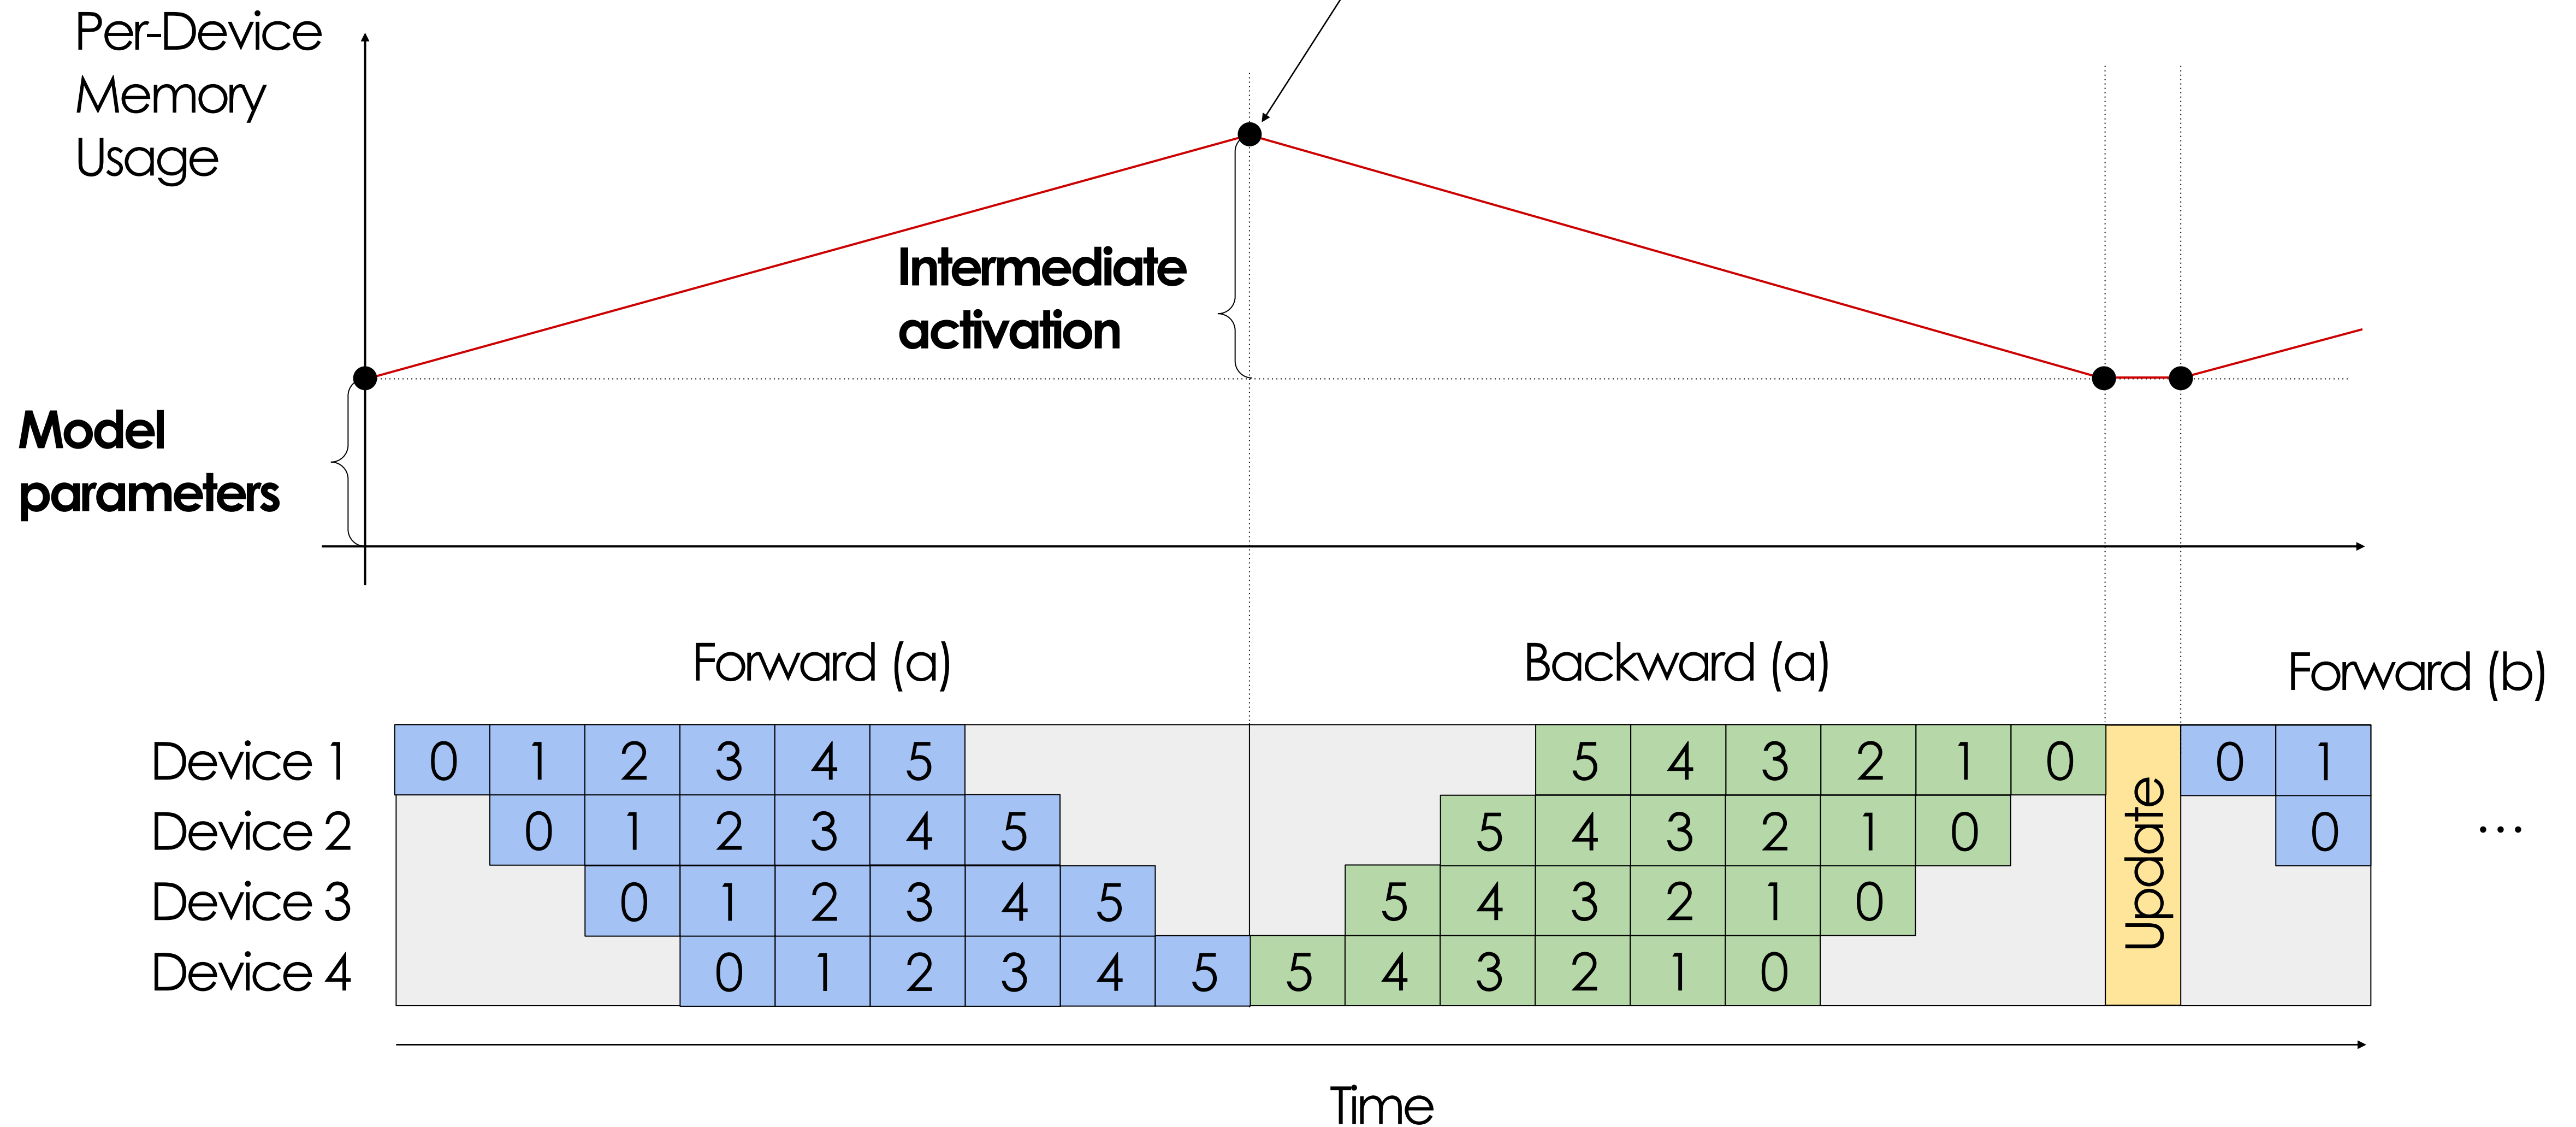
Example: Slice each input batch into 6 micro-batches:



Pipeline bubbles percentage = $(D - 1) / (D - 1 + N)$
with D devices and N micro-batches.


GPipe: Memory Usage

= Parameters + Activation \times #Micro-Batches




GPipe Schedule:

4
0



1F1B (1 Forward 1 Backward) Schedule:



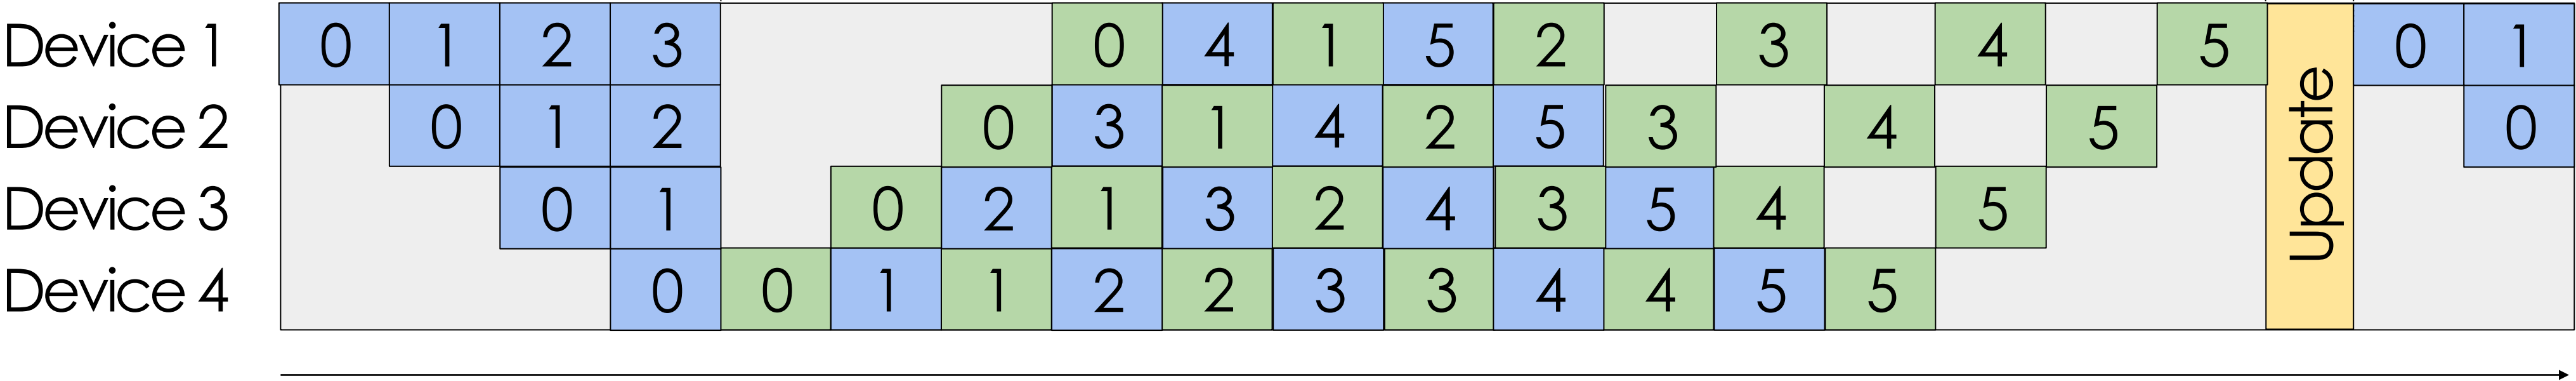
Same Latency

Perform backward as early as possible

1F1B Memory Usage

Maximum
per-device
memory
usage

= Parameters + Activation × ~~#Micro-Batches~~ #Devices



Time


...

Where We Are


- Deep Learning as Dataflow Graphs
- Auto-differentiation Libraries
 - Symbolic vs. Imperative
 - Static vs. Dynamic
- DL Parallelism
 - Inter-op parallelism
 - **Intra-op parallelism**

Recap: Intra-op and Inter-op

Strategy 1: Inter-operator Parallelism




Strategy 2: Intra-operator Parallelism



This section:

How to parallelize an **operator** ?
How to parallelize a **graph** ?




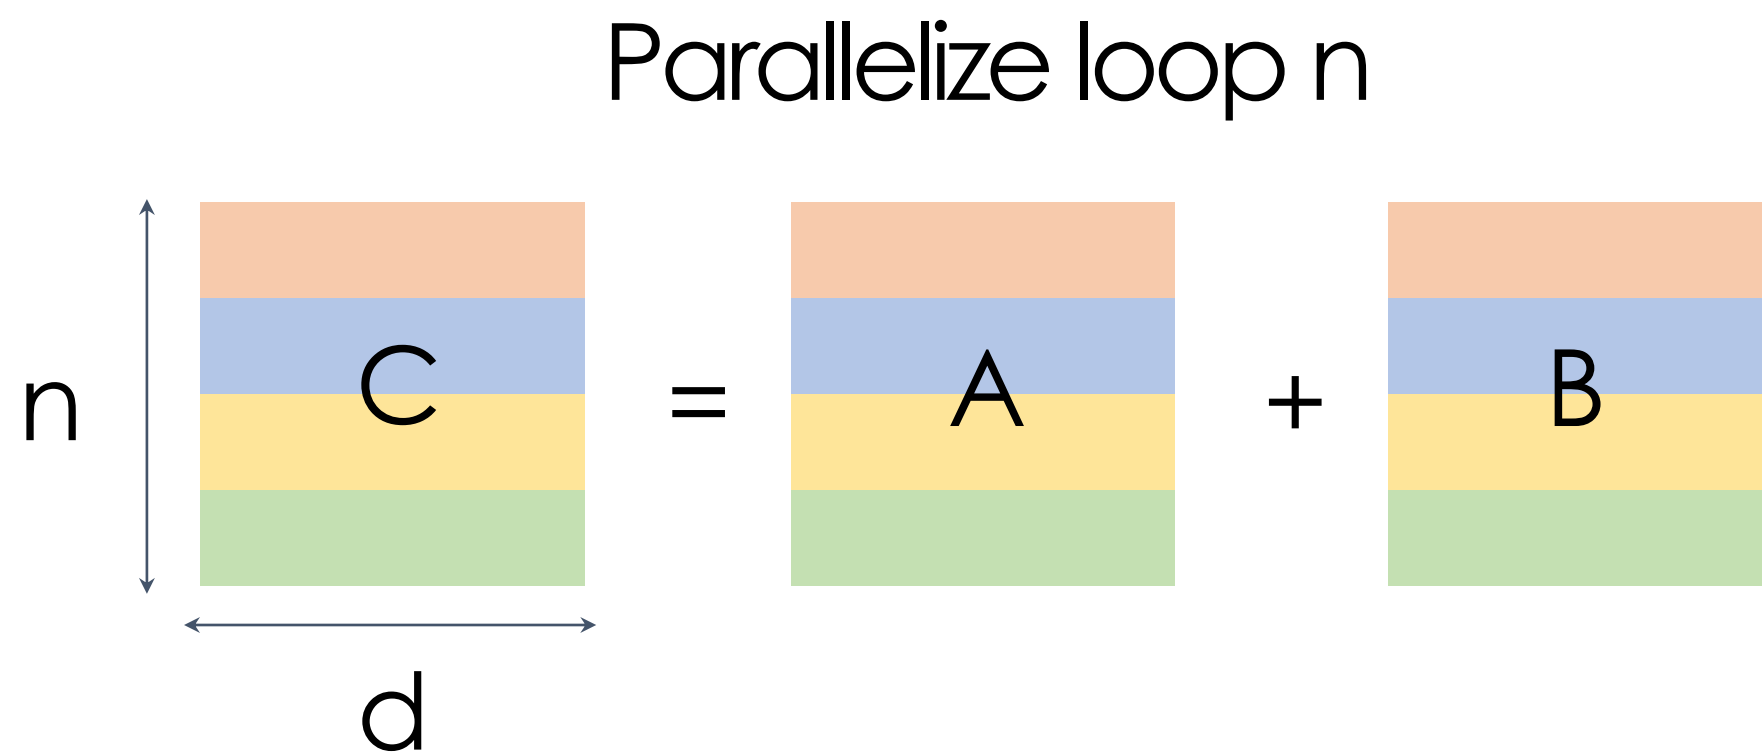
Parallelize One Operator

Element-wise operators

```
for n in range(0, N):  
  for d in range(0, D):  
    C[n,d] = A[n,d] + B[n,d]
```

No dependency on the two for-loops.
Can arbitrarily split the for-loops on different devices.

device 1 device 2 device 3 device 4



a lot of
other
variants
...

Parallelize One Operator


Matrix multiplication

```
for i in range(0, N):  
  for j in range(0, M):  
    for k in range(0, K):  
      C[i,j] = C[i,j] + A[i,k] x B[k,j]
```

No dependency on the two spatial for-loops.
Can arbitrarily split the for-loops on different devices.

Accumulation on this reduction loop.
Have to accumulate partial results if we split this for-loop

device 1 device 2 device 3 device 4 replicated



$$\begin{bmatrix} C_1 \\ C_2 \\ C_3 \\ C_4 \end{bmatrix} = \begin{bmatrix} A_1 \\ A_2 \\ A_3 \\ A_4 \end{bmatrix} \times B$$

Parallelize One Operator

Matrix multiplication

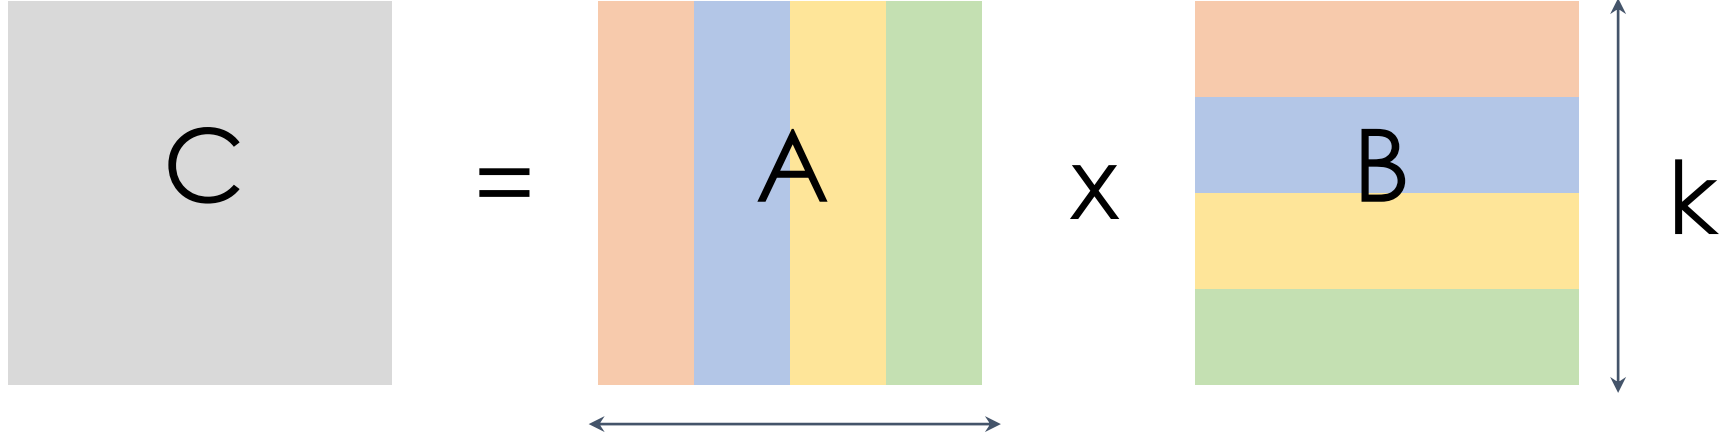
```
for i in range(0, N):  
  for j in range(0, M):  
    for k in range(0, K):  
      C[i,j] = C[i,j] + A[i,k] x B[k,j]
```

No dependency on the two spatial for-loops.
Can arbitrarily split the for-loops on different devices.

Accumulation on this reduction loop.
Have to accumulate partial results if we split this for-loop

device 1 device 2 device 3 device 4 replicate
d

Parallelize loop k


$$C = [A_1 \ A_2 \ A_3 \ A_4] \begin{bmatrix} B_1 \\ B_2 \\ B_3 \\ B_4 \end{bmatrix} = A_1 B_1 + A_2 B_2 + A_3 B_3 + A_4 B_4$$

(got by all-reduce) k

Parallelize One Operator

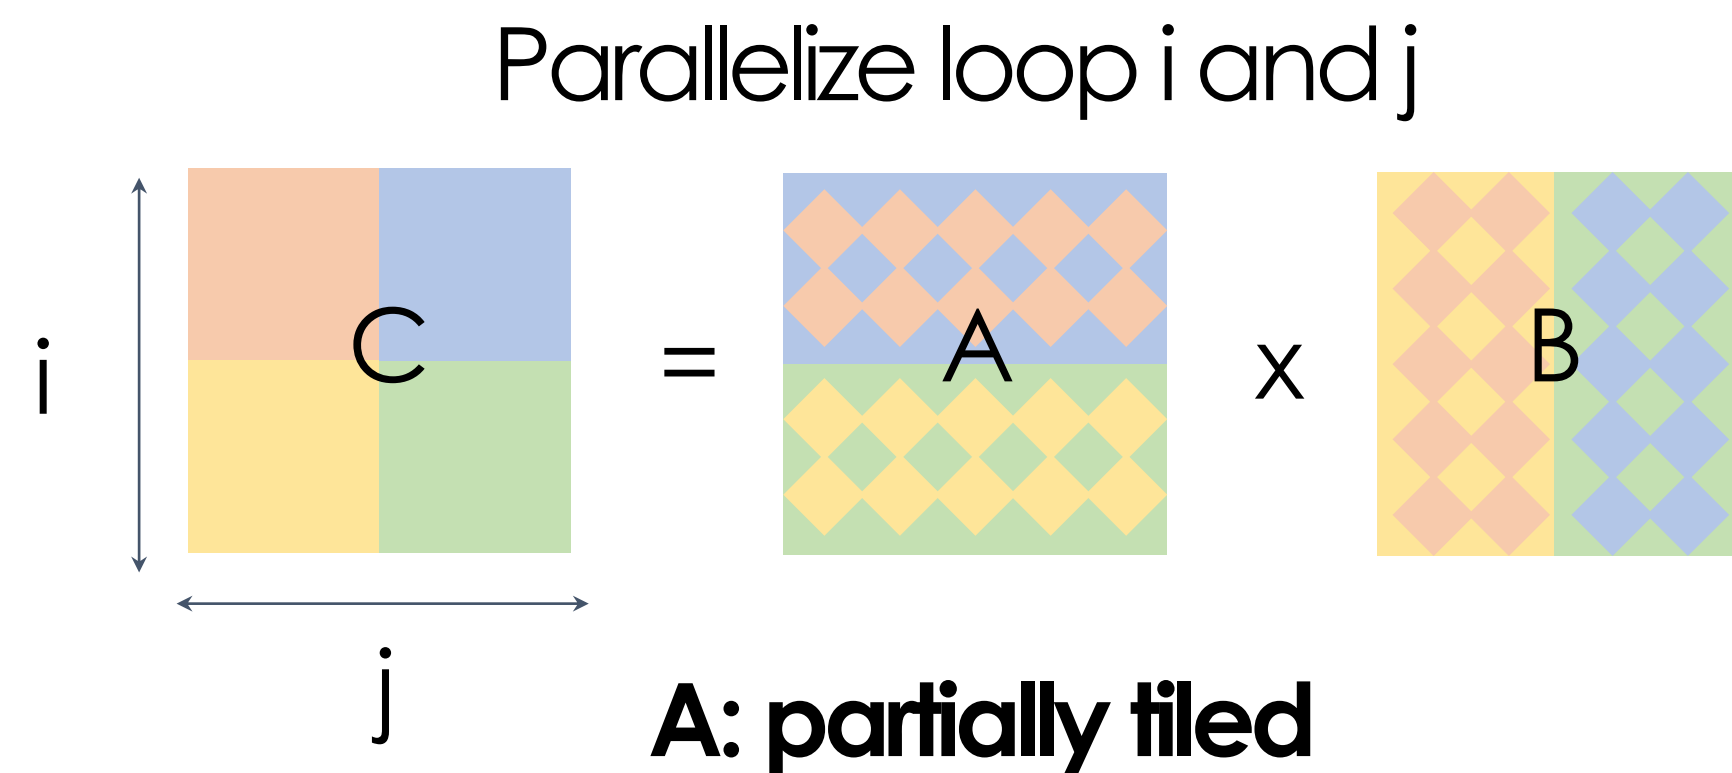
Matrix multiplication

```
for i in range(0, N):  
  for j in range(0, M):  
    for k in range(0, K):  
      C[i,j] = C[i,j] + A[i,k] x B[k,j]
```


No dependency on the two spatial for-loops.
Can arbitrarily split the for-loops on different devices.

Accumulation on this reduction loop.
Have to accumulate partial results if we split this for-loop

device 1 device 2 device 3 device 4



Device 1 and 2 hold a replicated tile
Device 3 and 4 hold a replicated tile



a lot of
other variants
...

Parallelize One Operator

2D Convolution

```
for n in range(0, N):  
  for co in range(0, CO):  
    for h in range(0, H):  
      for w in range(0, W):  
        for ci in range(0, CI):  
          for kh in range(0, KH):  
            for kw in range(0, KW):  
              C[n,co,h,w] += A[n,co,h+kh,w+kw] x B[kh,kw,co,ci]
```

Simple spatial loops. Can be arbitrarily split.

Stencil computation loops. Splitting these requires careful boundary handling.

Reduction loop. Need to accumulate partial results.

Reduction loops. But usually too small (≤ 5) for parallelization.

Simple case: Parallelize loop n , co , ci , then the parallelization strategies are almost the same as matmul's.


Complicated case: Parallelize loop h and w

Data Parallelism as A Case of Intra-op Parallelism

 Replicated  Row-partitioned  Column-partitioned


Matmul Parallelization Type 1

communication cost = 0




Matmul Parallelization Type 2

communication cost = all-reduce(c)




Forward Pass

Two “Type 1” matmuls: no communication



Backward Pass

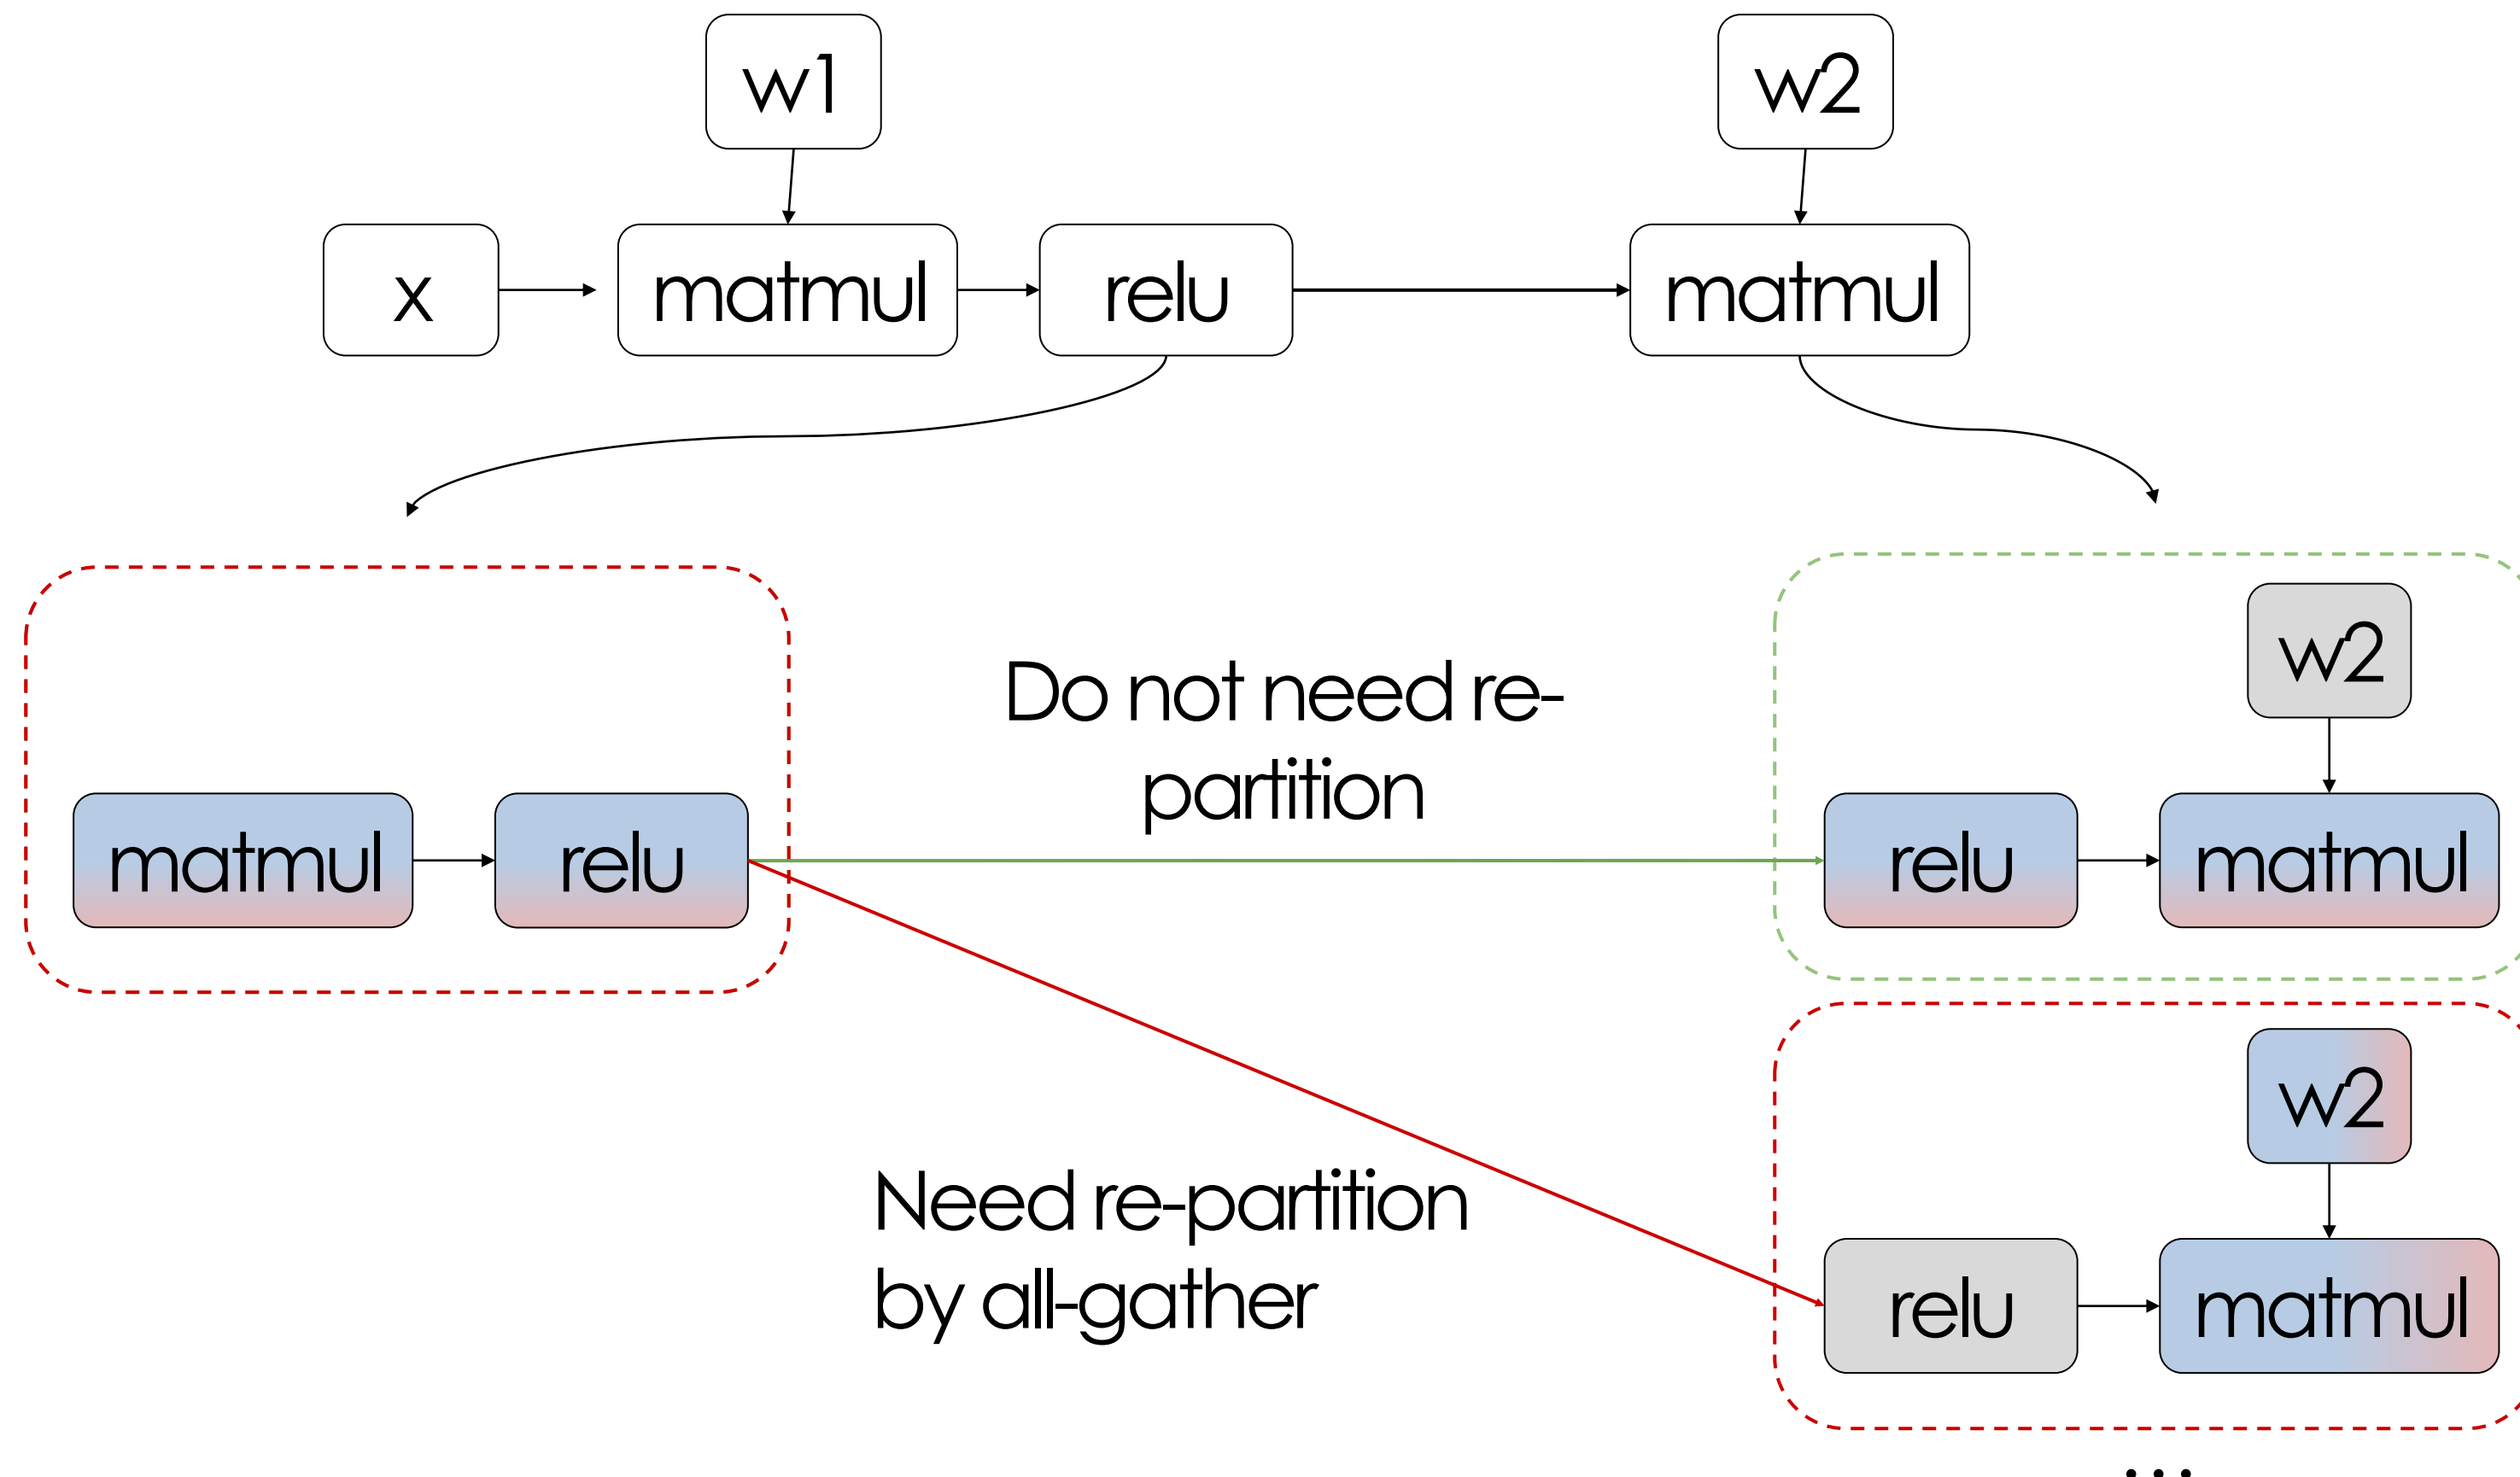
One “Type 1” matmul: no communication
Two “Type 2” matmuls: require all-reduce



Re-partition Communication Cost


Different operators' parallelization strategies require different partition format of the same tensor

■ Replicated ■ Row-partitioned ■ Column-partitioned



Re-partition Communication Cost


Different operators' parallelization strategies require different partition format of the same tensor



Parallelize All Operators in a Graph

Problem

Pick a parallel strategy
of each operator



Minimize **Node costs** (computation + communication) + **Edge costs** (re-partition communication)

Solution


Manual design
Randomized search
Dynamic programming
Integer linear programming

Important Projects

- Model-specific Intra-op Parallel Strategies
 - **AlexNet**
 - **Megatron-LM**
 - GShard MoE
- Systems for Intra-op Parallelism
 - ZeRO
 - Mesh-Tensorflow
 - GSPMD
 - Tofu
 - FlexFlow

AlexNet

Result: increase top-1 accuracy by 1.7%

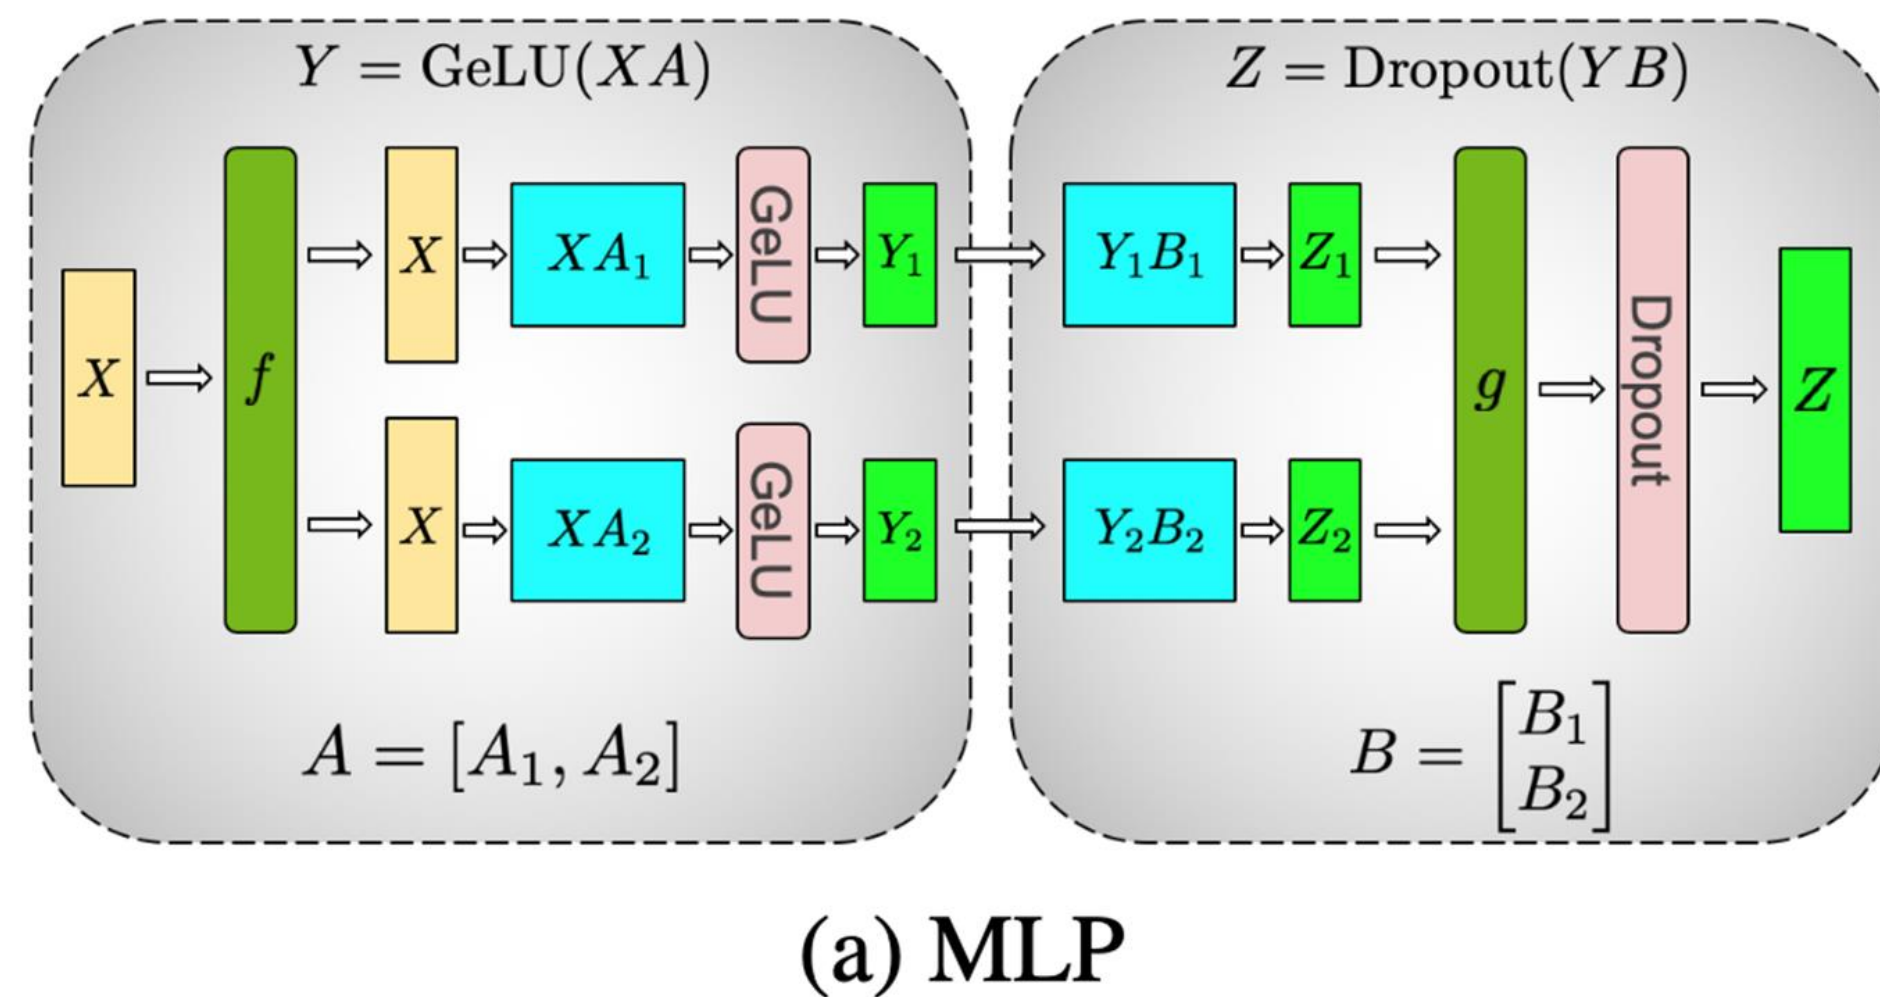


Assign a group convolution layer to 2 GPUs

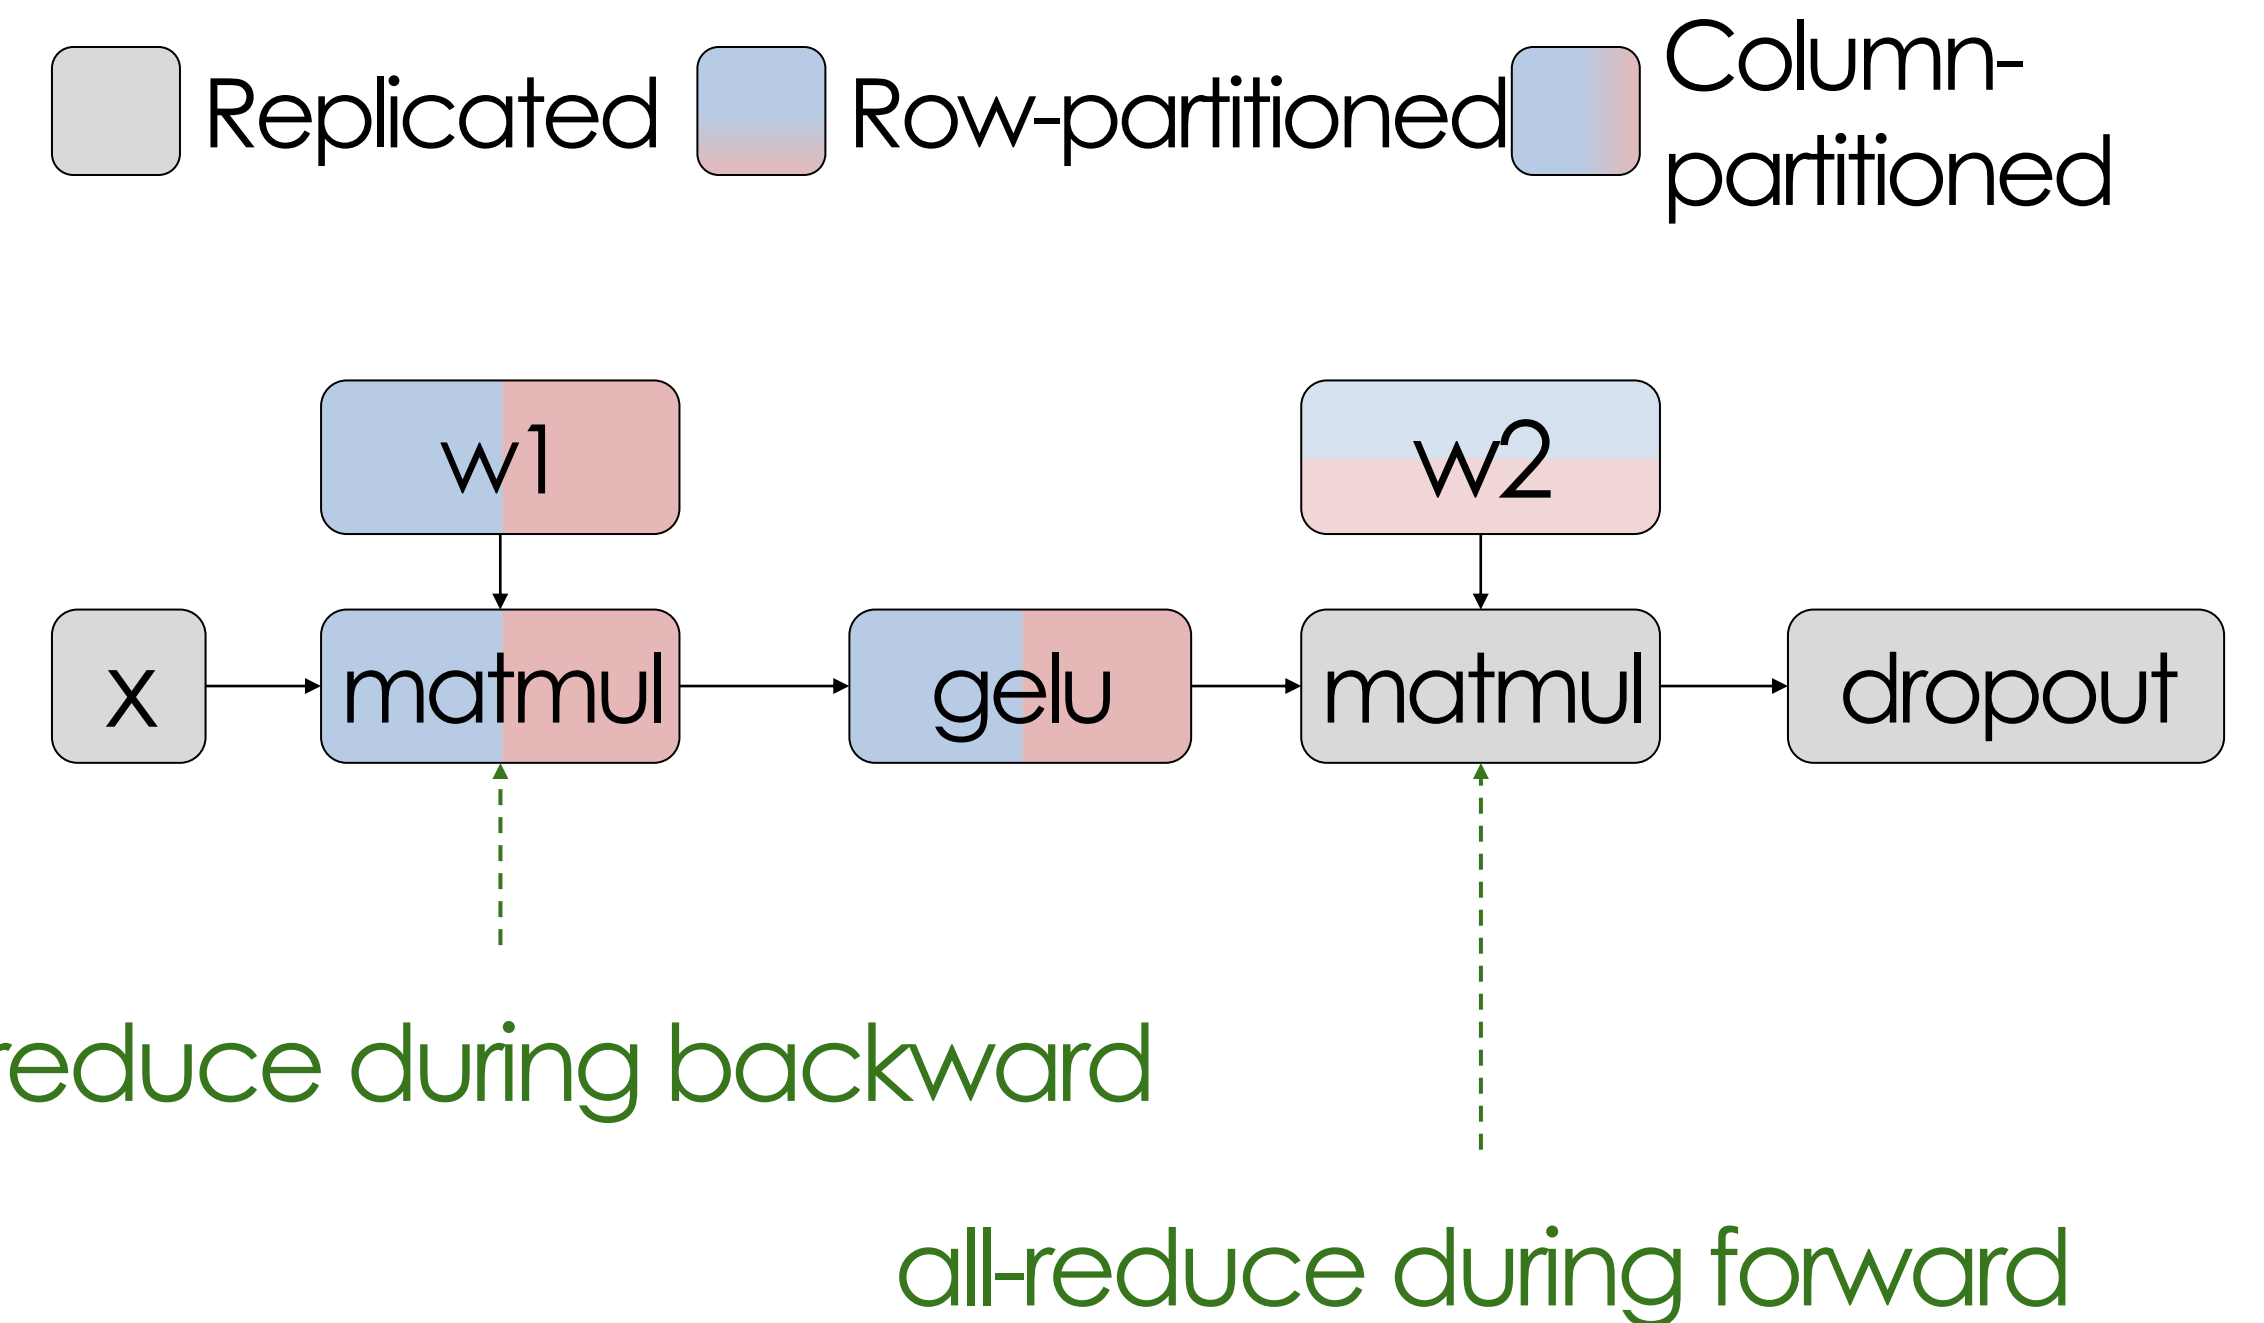
Megaton-LM

Result: a large language model with 8.3B parameters that outperforms SOTA results

Figure 3 from the paper:
How to partition the MLP in the transformer.

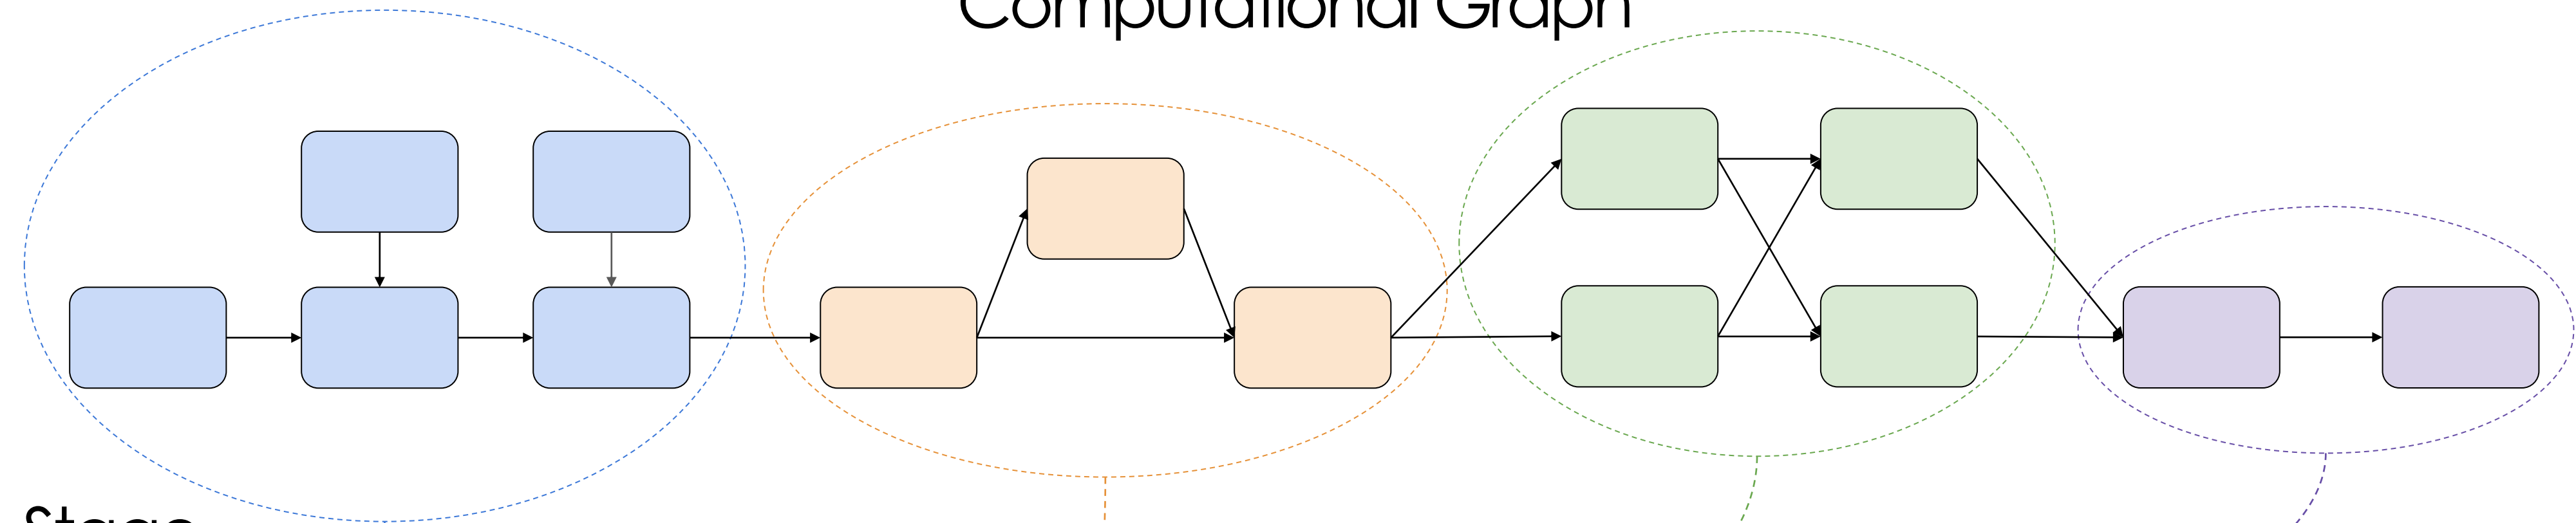


Illustrated with the notations in this tutorial



Combine Intra-op Parallelism and Inter-op Parallelism





Computational Graph



Stage

Device
Mesh

Intra-op Parallelism



Inter-op Parallelism

Intra-operator Parallelism Summary

- . We can parallelize a single operator by exploiting its internal parallelism
- . To do this for a whole computational graph, we need to choose strategies for all nodes in the graph to minimize the communication cost
- . Intra-op and inter-op can be combined